



SCHOOL OF COMPUTATION,
INFORMATION AND TECHNOLOGY —
INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

**Combatting the Precision Loss of Partial
Contexts in Abstract Interpretation**

Felix Sebastian Kraye



SCHOOL OF COMPUTATION,
INFORMATION AND TECHNOLOGY —
INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

**Combatting the Precision Loss of Partial
Contexts in Abstract Interpretation**

**Bekämpfung des Präzisionsverlustes durch
partielle Kontexte in Abstrakter
Interpretation**

Author:	Felix Sebastian Kraye
Supervisor:	Prof. Dr. Helmut Seidl
Advisor:	Michael Schwarz
Submission Date:	15th of February 2023

I confirm that this bachelor's thesis is my own work and I have documented all sources and material used.

Munich, 15th of February 2023

Felix Sebastian Kraye

Acknowledgments

Abstract

Contents

Acknowledgments	iii
Abstract	iv
1 Introduction	1
2 Background	2
2.1 Static Analysis	2
2.1.1 Flow sensitive analysis	3
2.1.2 Constraint systems	4
2.1.3 Interprocedural analysis	4
2.1.4 Context sensitivity	6
2.1.5 Partial context sensitivity	7
2.1.6 Precision loss	8
3 Combatting Precision Loss	10
3.1 Formal description	10
3.1.1 Taint analysis	10
3.1.2 Improving the values-of-variables analysis	11
3.2 Implementation	12
3.2.1 Taint analysis	13
3.2.2 Benefiting other analyses	15
4 Evaluation	19
4.1 Testing	19
4.2 Benchmarking	19
5 Conclusion	20
Abbreviations	21
List of Figures	22
List of Tables	23

Bibliography

24

1 Introduction

WIP:

- introduce GOBLINT
- show problem on a small example

Related work

Structure First we will introduce the basics of static analysis. This will go by introducing constraint systems and how these are used to gain information about the program statically. It will be accompanied by an example of a value-of-variables analysis acting on a toy language we will use for examples in this thesis. This will be extended to an interprocedural approach where partial context sensitivity will be introduced. Here the source of the precision loss will be pointed out. We then will propose an approach to combat this precision loss. The approach will first be introduced theoretically, after which we also present the challenges and results of implementing it in the GOBLINT analyzer. To give an evaluation to the proposed approach, a benchmark of the implementation will be performed and inspected. Our conclusions are presented in the last chapter.

2 Background

2.1 Static Analysis

Static analysis is defined by Rival [RY20] as "[...]an automatic technique that approximates in a conservative manner semantic properties of programs before their execution". This means that the program is analyzed just by the given source code without execution. The goal is to prove certain properties about the program in a "sound" manner, i.e., any property that is proven to hold actually does hold. However, from failing to prove a property one cannot conclude that the given property does not hold.

In order to prove properties, e.g. finding that a program does not contain races or identifying dead code, we need to gain information about the program. This is done by performing various kinds of analyses. We will focus on flow sensitive analyses from now on, i.e., analyses which find properties of the program dependent on the location within it. We will introduce a syntax to formalize flow sensitive analyses in the following sections. This formalization approach is heavily based on [ASV12].

// TODO: talk about Abstract Interpretation.

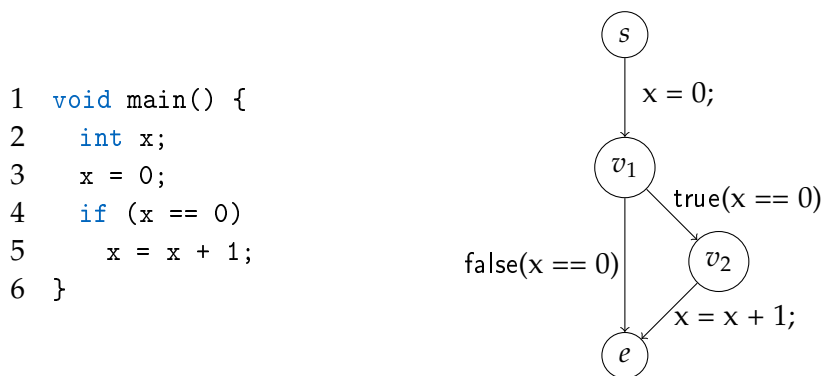


Figure 2.1: Example program (left) and corresponding CFG (right)

2.1.1 Flow sensitive analysis

As noted above flow sensitive analyses find properties of the program dependent on the point within the program. Expressed differently this means a flow sensitive analysis will find an overapproximation of states the program may be in for any given point within the program or "program point". This state can describe many things dependent on the analysis performed.

First let us define what a program point is: Consider a Control flow Graph (CFG), where nodes represent points between instructions within the program. Edges are labeled with instructions or checks (from now on collectively called "actions") and describe the transitions between these points (see example Figure 2.1). Then any node on this CFG would be what we call a program point.

Concretely let N be the set of all program points. Furthermore, let \mathbb{D} be a Domain containing abstract states describing concrete states of the program. This means that some $d \in \mathbb{D}$ can describe many states the program can be in.

Then an analysis is expected to find a mapping $\eta : N \rightarrow \mathbb{D}$ which maps program points to abstract states describing that location within the program, i.e., for $[n] \in N$, $\eta [n]$ should be an abstract state describing all possible states (and possibly more) the program can be in at program point $[n]$.

As an example we will introduce a values-of-variables analysis for integers. This analysis finds a mapping from a set of program variables X to abstractions of their possible values at any given program point. Our toy language will support global variables (globals) as well as local variables (locals). The global variables can be accessed and changed by any procedure, while local ones are only visible to the procedure in which it was declared and can only be accessed and changed by this procedure. Therefore, our set X of variables is the disjoint union of globals G and locals L : $X = G \uplus L$. In the scope of this thesis we will focus on abstracting integer values by sets of integers. Thereby the goal of our values-of-variables analysis is to find a mapping $X \rightarrow 2^{\mathbb{N}}$ for each program point.

Combining this with the considerations from above, we get that the Domain \mathbb{D}_v for the values-of-variables analysis should be $\mathbb{D}_v = X \rightarrow 2^{\mathbb{N}}$. Finally, the resulting $\eta_v : N \rightarrow \mathbb{D}_v$ for this analysis describes a mapping $\eta_v [n]$ for some program point $[n] \in N$, where $\eta_v [n] x$ is a set containing all values $x \in X$ may possibly hold at $[n]$. From this we can conclude that x cannot hold any value outside $\eta_v [n] x$ at program point $[n]$.

2.1.2 Constraint systems

We now formulate a way in which we can describe an analysis in the form of constraints. For this we need a partial ordering \sqsubseteq on the domain \mathbb{D} .

Then we create a system of constraints which can be solved for a solution. Consider the edges (u, A, v) of the CFG, where each edge denotes a transition from program point $[u]$ to program point $[v]$ via the action A . Now let each of these edges give rise to a constraint

$$\eta [v] \sqsubseteq \llbracket A \rrbracket^\# (\eta [u])$$

where $\llbracket A \rrbracket^\#$ denotes the abstract effect of the action A defining our analysis. In addition, we need a start state. This is given by $\text{init}^\# : \mathbb{D}$ which is defined depending on the analysis. This gives rise to the start constraint $\eta [s] \sqsubseteq \text{init}^\#$ for the starting point of the program $[s] \in N$.

We will show these ideas with our example of the values-of-variables analysis: Let us define the partial ordering \sqsubseteq_v that is necessary for building the constraints. We will do this by saying that a mapping $M_1 \in \mathbb{D}_v$ is ordered above another mapping M_2 iff for every variable the set it is mapped to in M_1 is a superset of the one the variable is mapped to in M_2 . Formulated formally this is:

$$M_1, M_2 \in \mathbb{D}_v : M_1 \sqsubseteq_v M_2 \iff \forall x \in X : M_1 x \subseteq M_2 x$$

Next we define the start state $\text{init}^\# = M_\top$ for this domain as the mapping that maps every variable to the full set of integers \mathbb{N} , i.e., $\forall x \in X : M_\top x = \mathbb{N}$. This is because we assume variables to be randomly initialized in our toy language.

It remains to define the abstract effect of actions $\llbracket A \rrbracket_v^\#$ for our values-of-variables analysis. We will just show the effect of a simple variable assignment:

$$\llbracket x = y; \rrbracket_v^\# M = M \oplus \{x \mapsto (M y)\}$$

where $M \oplus \{x \mapsto s\}$ denotes that the mapping M is updated such that x will be mapped to the set s . A full definition of abstract effects of a values-of-variables analysis can be found at </ / TODO >.

2.1.3 Interprocedural analysis

So far we only have defined how a program without procedure calls is analyzed. Now we want to introduce procedure calls of the form $f()$. For simplicity, we will only consider argumentless procedure calls without a return value. Arguments and return

values can be simulated by using global variables.

Since a call has its own set of local variables to work with and a call stack can contain multiple of the same procedure (e.g. for recursion), we will analyze procedures in their own environment. However, we need to consider global variables and how the procedure affects these.

The idea is to give procedures their own starting states and analyze them similarly as we have done before. The final state of the called procedure is then used to be combined back into the state of the caller before the call. Formalized for an edge $(u, f();, v)$ this looks as follows:

$$\begin{aligned}\eta[s_f] &\sqsupseteq \text{enter}^\#(\eta[u]) \\ \eta[v] &\sqsupseteq \text{combine}^\#((\eta[u]), (\eta[e_f]))\end{aligned}$$

where $[s_f]$ and $[e_f]$ are the start and end node of the CFG for procedure $f()$. The functions $\text{combine}^\# : \mathbb{D} \times \mathbb{D} \rightarrow \mathbb{D}$ and $\text{enter}^\# : \mathbb{D} \rightarrow \mathbb{D}$ are defined by the analysis. $\text{enter}^\#$ handles computing the start state for the procedure $f()$, while $\text{combine}^\#$ describes in what way the caller state and the end state of the callee are merged after the call.

It is worth mentioning at this point that even though a procedure can be called from multiple points within the program we still only analyze the procedure once. For n procedure calls $(u_n, f();, v_n)$ we get n constraints for $[s_f]$: $\eta[s_f] \sqsupseteq \text{enter}^\#(\eta[u_n])$. We can express this differently in a single constraint as follows:

$$\eta[s_f] \sqsupseteq \bigsqcup \{d \mid \exists (u_n, f();, v_n) \in \text{Edges} : \text{enter}^\#(\eta[u_n]) = d\}$$

where \bigsqcup is the least upper bound, i.e., the least $d \in \mathbb{D}$ according to the ordering \sqsubseteq that is ordered above all of its argument elements.

For our values-of-variables analysis we will show how $\text{enter}_v^\#$ and $\text{combine}_v^\#$ are defined. We need to take global variables into account when computing the start state and combining after the call. Therefore, we define the two functions as follows:

$$\begin{aligned}\text{enter}_v^\# M &= M|_{\text{Globals}} \oplus \{x \mapsto \mathbb{N} \mid \forall x \in X\}|_{\text{Locals}_{ce}} \\ \text{combine}_v^\#(M_{cr}, M_{ce}) &= M_{cr}|_{\text{Locals}_{cr}} \oplus M_{ce}|_{\text{Globals}}\end{aligned}$$

where $M|_{\text{Locals}}$ and $M|_{\text{Globals}}$ refers to the mapping M restricted to only the local or global variables respectively. Note that Locals_{ce} refers to the locals of the callee while Locals_{cr} refers to the locals of the caller.

To explain these two function let us first look at $\text{enter}_v^\#$. This function takes the part of the mapping from the caller that contains information about global variables and adds the information of uninitialized local variables used in the procedure to the state. For $\text{combine}_v^\#$ the local part from the callee is kept, but it is updated with the global part of

the callee return state as this contains the updated information about global variables after the procedure call.

2.1.4 Context sensitivity

In the previous chapter we approached the analysis of procedures by analyzing them only once with an abstract start state describing all possible concrete states the procedure could start with. We call this behavior "context insensitive" as the procedure is analyzed without differentiating between different contexts it is called in.

This is not very precise as we will exemplify by applying the values-of-variables analysis to the program in Figure 2.2. We ignore the marked lines of the program for now.

The procedure `incr()`; is called twice: Once with $a = 1$ in Line 10 and once with $a = -3$ in Line 13. This leads to two constraints for node $[s_{incr}]$:

$$\eta_v [s_{incr}] \sqsupseteq_v \text{enter}_v^\# \eta_v [v_2] = \{a \rightarrow \{1\}\}$$

$$\eta_v [s_{incr}] \sqsupseteq_v \text{enter}_v^\# \eta_v [v_5] = \{a \rightarrow \{-3\}\}$$

leading to $\eta_v [s_{incr}] = \{a \rightarrow \{-3, 1\}\}$. At the end point of the call the state will be $\eta_v [e_{incr}] = \{a \rightarrow \{-2, 2\}\}$, which is then combined back into the states of nodes in the main procedure for node $[v_6]$ and propagated up to the `assert(a < 0)`; in Line 14. The result of this assertion cannot be determined by the analysis even though it is easy for humans to see that it should hold.

This could have been avoided, if the procedure was analyzed twice, once with each starting state. To achieve this we will need to perform some modifications on our current approach: Instead of searching a mapping $\eta : N \rightarrow \mathbb{ID}$ we from now on seek $\eta : (N \times \mathbb{ID}) \rightarrow \mathbb{ID}$. This allows us to have different states for the same program point. We call the second part of $N \times \mathbb{ID}$ "context". For now this context will be the same as the starting state of the current procedure. Therefore, we need to adjust the notion of $\text{enter}^\#$ and $\text{combine}^\#$:

$$\eta [s_f, \text{enter}^\# (\eta [u, d])] \sqsupseteq \text{enter}^\# (\eta [u, d])$$

$$\eta [v, d] \sqsupseteq \text{combine}^\# ((\eta [u, d]), (\eta [e_f, \text{enter}^\# (\eta [u, d])]))$$

The main procedure will always be analyzed just once as in our toy language it is only called initially. The context for its nodes can be chosen arbitrarily.

There are no changes we need to perform on the values-of-variables analysis to make it

context-sensitive. Solely the changes to the general analysis framework above suffice. Applying this analysis to the example in Figure 2.2 would lead to the procedure `incr()` being analyzed twice with different contexts, assuming we still ignore the marked lines. This leads to the following two entry constraints for different unknowns of the constraint system:

$$\begin{aligned}\eta_v [s_{incr}, \{a \rightarrow \{1\}\}] &\sqsupseteq_v \{a \rightarrow \{1\}\} \\ \eta_v [s_{incr}, \{a \rightarrow \{-3\}\}] &\sqsupseteq_v \{a \rightarrow \{-3\}\}\end{aligned}$$

For node $[v_6]$ only the state $\eta_v [e_{incr}, \{a \rightarrow \{-3\}\}] = \{a \rightarrow \{-2\}\}$ is combined into the caller state before the call. With this information we can safely say that the assertion in the following Line 14 will hold.

2.1.5 Partial context sensitivity

While the context-sensitive approach from the previous section might be very precise, it can be quite costly in terms of computation time. To reach a middle ground between a context insensitive and a fully context-sensitive analysis, one could propose an approach where only a part of the domain \mathbb{ID} is used as a context to differentiate function calls. Let $\mathbb{ID} = \mathbb{ID}_{\text{ctx}} \times \mathbb{ID}_{\text{rem}}$. The domain is now a product of one part used for contexts \mathbb{ID}_{ctx} and the remaining part \mathbb{ID}_{rem} .

Again we need to change the definition of the mapping we want to compute: We now have $\eta : (N \times \mathbb{ID}_{\text{ctx}}) \rightarrow \mathbb{ID}$. Additionally, `enter#` and `combine#` are changed as follows:

$$\begin{aligned}\eta [s_f, \langle \text{enter}^\# (\eta [u, c]) \rangle_1] &\sqsupseteq \text{enter}^\# (\eta [u, c]) \\ \eta [v, c] &\sqsupseteq \text{combine}^\# ((\eta [u, c]), (\eta [e_f, \langle \text{enter}^\# (\eta [u, c]) \rangle_1]))\end{aligned}$$

where $\langle d \rangle_1$ extracts the first part of a tuple $d \in (\mathbb{ID}_{\text{ctx}} \times \mathbb{ID}_{\text{rem}})$. This formalization results in multiple constraint for a single starting variable with context $[s_f, c]$. We can alternatively formulate this as

$$\eta [s_f, c] \sqsupseteq \bigsqcup \{(c, d) \mid \exists (u_n, f(), v_n) \in \text{Edges} : \text{enter}^\# (\eta [u_n, c_n]) = (c, d)\}$$

i.e., the constraint for the variable $[s_f, c]$ is the least upper bound of all entry states (c, d) for some call of f , which have the same context c as the variable. Or expressed differently, all states (c, d) computed by `enter#` for f are grouped by the context c , where each group is joined by \bigsqcup to produce a constraint for a starting variable $[s_f, c]$ with the respective context.

With this formal model we have the option to perform an analysis completely context sensitively ($\mathbb{ID} = \mathbb{ID}_{\text{ctx}} \times (\bullet)$), completely context insensitively ($\mathbb{ID} = (\bullet) \times \mathbb{ID}_{\text{rem}}$) or anything in between, where $\{(\bullet)\}$ is the "unit domain" which contains exactly one

element with the trivial ordering $(\bullet) \sqsubseteq (\bullet)$.

We have to note here that there are some severe issues with the approach for (partial) context-sensitive analysis described in this thesis: The resulting system of constraints may not be finite and some variables in the constraint system may depend on an infinite number of other variables. This can result in a very hard or even impossible to solve constraint system. While we will stick with this formal approach for this thesis for the sake of simplicity, an escape route to this issue is described in [ASV12].

2.1.6 Precision loss

The main source of the precision loss of context-insensitive or partially context-sensitive analysis is the join over all states with the same context. Consider a procedure f that has no effect, i.e., $s_f = e_f$. Even for this procedure, the $\text{combine}^\#$ function receives the less precise result of the join to combine it with the caller state. In this case it would be more precise, if it directly used the result from $\text{enter}^\#$ as the callee state for combining. Even for procedures that do change the state, there might be some parts of the state which are untouched by the call. If we can identify these untouched parts, we could reduce the precision loss experienced by using partial contexts.

Let us clarify the source of the precision loss mentioned above with an example: For this we once again consider the example program Figure 2.2. This time we take the marked lines into account. When the program is analyzed context insensitively, not only will the state at the start node represent two possible values for the global variable a , but also for c . The state will therefore be

$$\eta_v [s_{\text{incr}}, (\bullet)] = \{a \rightarrow \{-3, 1\}, c \rightarrow \{-10, 10\}\}$$

Even though the variable c is never changed within $\text{incr}()$, the mapping $c \rightarrow \{-10, 10\}$ is still copied into the caller state when combining the states for node $[v_5, (\bullet)]$. Therefore, the assertion in Line 15 cannot be determined to hold solely with the information gained by context insensitive values-of-variables analysis. This precision loss could easily be avoided if we had some idea which global variables are definitely not changed by a procedure call.




Figure 2.2: Example program (left) and corresponding CFGs for main (middle) and incr (right)

3 Combatting Precision Loss

In this chapter we will describe our approach to reduce the precision loss described in Subsection 2.1.6. We will first use the syntax for flow sensitive analyses from Chapter 2 to formally define the idea. After that we explain the concrete implementation of the approach into the GOBLINT analyzer.

3.1 Formal description

3.1.1 Taint analysis

The basic idea to combat the precision loss is to track for each procedure which variables have been written or have possibly been altered in some other way. This information is then used in the values-of-variables analysis when combining the abstract state from the caller with the abstract return state given by the callee at the end of the procedure. In the following we will call a variable that has been written or altered in the current procedure context "tainted". Therefore, we introduce a new taint analysis tracking which variables have been tainted within the context of the current procedure. It is worth mentioning that our notion of taintedness is related but different from other uses of this concept.

Let us now formulate the syntax for our taint analysis: Since we want to find a collection of tainted variables per program point, a suitable domain for this analysis is the powerset of the set of variables X ordered by the subset relation:

$$\mathbb{D}_t = 2^X \text{ with } \sqsubseteq_t = \subseteq$$

From that follows that we seek to compute a mapping from program points to sets of variables, i.e., $\eta_t : N \rightarrow \mathbb{D}_t$. To interpret this with the goal of our taint analysis in mind, we note that $\eta_t[n, (\bullet)] = T$ will denote that T is the set of possibly tainted variables at program point $[n]$. Expressed differently this means that for any variable $x \in T$ we cannot exclude that this variable was altered between the start of the procedure $[n]$ is in up until the program point $[n]$.

It remains to define $\text{init}_t^\#$, $\text{enter}_t^\#$ and $\text{combine}_t^\#$ as well as the abstract effects of actions $\llbracket A \rrbracket^\#$. Recall that the notion of a "tainted" variable is defined in relation to the current

procedure. This means we want to start fresh whenever we enter a procedure and start without any variable being initially tainted. It is worth pointing out that the entry to a procedure call does not depend on the state where it is called. Therefore, we design our analysis to be context-insensitive. With these considerations we define $textsfer_{\tau}^{\#}$ and $init_{\tau}^{\#}$ as follows:

$$enter_{\tau}^{\#} T = init_{\tau}^{\#} = \emptyset$$

It is worth pointing out here that the function $enter_{\tau}^{\#} T$ is always equal to the empty set irregardless of its argument T . Therefore, it computes the same function context for each call of a certain procedure making our taint analysis inherently context insensitive. When combining the caller state with the returned callee state, we note that anything that we need to keep the tainted set from before the call, as a tainted variable can get never get "untainted" again, no matter what the procedure does. In addition to that we will add the set returned by the callee, as anything tainted in the call needs to be considered tainted in the caller as well. This is because we want to know which variables have been altered in a procedure call, no matter if the tainting happened within the procedure itself or within a procedure called by the procedure. This leaves us with the following equation for the $combine_{\tau}^{\#}$ function:

$$combine_{\tau}^{\#} (T_{cr}, T_{ce}) = T_{cr} \cup (T_{ce} \setminus Local_{ce})$$

Note that we removed the callee local variables $Local_{ce}$ because these are not accessible by the caller and all of its callers anyway, so it is not useful to keep track of them. Lastly we define the abstract effects of actions. Most of these (including checks) do not do anything besides propagating through the state from before. The only major exception is a variable assignment. For these we note that this specific variable, which the value is assigned to is added to the tainted set. This is independent of the expression that evaluates to the assigned value, as we are only interested in the fact that the variable on the left of the assignment is altered. This leaves us with the following abstract effects of actions:

$$\llbracket A \rrbracket^{\#} T = \begin{cases} T \cup \{x\} & \text{if } A \equiv (x = e;) \\ T & \text{else} \end{cases}$$

where e is any arbitrary expression.

This concludes our definition of the taint analysis. In the following chapter we will see how this information helps us to improve the values-of-variables analysis.

3.1.2 Improving the values-of-variables analysis

Recall the source of the precision loss we want to reduce. This happened when a global variable was updated with a less precise value after a procedure call even though this

specific variable was not changed by the call.

Thanks to the taint analysis we defined above, we now do have the information which variables can be altered by a procedure $f()$ and which surely stay untouched. These are exactly those variables which are not in the tainted set of the end node $[e_f]$ for that procedure.

With this insight we can now update the $\text{combine}_v^\#$ function of our values-of-variables analysis as follows:

$$\text{combine}_v^\#(M_{cr}, M_{ce}) = M_{cr}|_{\text{Locals}_{cr} \cup (\text{Globals} \setminus T_{ce})} \oplus M_{ce}|_{\text{Globals} \cap T_{ce}}$$

where for an edge $(u, f();, v)$ we have $T_{ce} = \eta_t[e_f]$.

Similar to before the $\text{combine}_v^\#$ function takes the caller mapping, restricts it to a subset of caller reachable variables and updates this mapping with the callee mapping restricted to the rest of caller reachable variables. In other words, the caller reachable variables are partitioned into two sets such that one subset is taken from the caller state while the other one is taken from the callee state. Before this change the partitioning was done strictly in such a way that the local variables were taken from the caller state and all global variables from the callee state. After this change, the global variables that are not tainted by the callee are also taken from the caller state and not from the callee anymore. Thereby the precision loss for untainted variables is eliminated.

One might wonder if this change could lead to a case, where the callee state has a more precise value for a variable that is discarded because this variable is not in the tainted set. Concretely this situation would be described by

$$\exists \text{Edge } (u, f();, v), x \in \text{Globals} : x \notin \eta_t[e_f] \wedge (\eta_v[e_f] x \subset \eta_v[u] x)$$

From $x \notin \eta_t[e_f]$ we know that x has not been altered in the procedure $f()$ since the node $[s_f]$, and therefore it holds that

$$\eta_v[e_f] x = \eta_v[s_f] x$$

By the definitions of \sqsubseteq_v and $\text{enter}_v^\#$ we get:

$$\eta_v[s_f] x \supseteq (\text{enter}_v^\#(\eta_v[u])) x = \eta_v[u] x$$

Therefore, $\eta_v[e_f] x \supseteq \eta_v[u] x$ which is a contradiction to the proposed case which we can therefore exclude.

3.2 Implementation

We will quickly introduce the GOBLINT analyzer and its structure before we explain the process of implementing the proposed taint analysis as well as its usage to improve

other analyses. The core functionality of GOBLINT is to statically analyze C programs using an approach similar to the one described in Chapter 2. This generally works as follows: After the C input file is preprocessed, a CFG is generated. This is then used together with the specifications of various analyses to generate a constraint system. GOBLINT solves this constraint system and produces different kinds of outputs to the user according to the solution (e.g. notifications, warnings or a visualization of the full solution).

It is worth mentioning that GOBLINT can perform multiple analyses on a program at the same time. For this a compound domain is built that is a tuple of all the domains of the analyses to be performed. To generate constraints, all activated analyses are taken into account where the specification of each analysis acts on its corresponding part of the compound domain. Information can be transferred between the different analyses via a system called "queries".

Figure 3.1 shows the inner structure of the analyzer. We can see that GOBLINT provides parametrized domains which can be used in the specifications of the analyses. It is also shown that multiple analyses are then combined into one MCP that is then used with the CFG to generate constraints which are solved.

For a deeper insight into the inner workings of GOBLINT refer to [Api14].

3.2.1 Taint analysis

To define an analysis the GOBLINT analyzer provides an interface, where the relevant parts can be seen in Figure 3.2. This interface requires two modules `D` and `C` which define the domain and the context-sensitive part of the domain. After that some functions are required:

- `name` to uniquely refer to an analysis.
- `startstate` to define the state used when entering the analysis (similar to `init#`).
- `query` to implement the query system of GOBLINT. This allows an analysis to broadcast information to be used in analyses.
- *Transfer functions* which define the abstract effects of actions (similar to $\llbracket A \rrbracket^{\#}$)
- *Functions for interprocedural analysis*
- *Function for analysis of multithreaded programs*

For our taint analysis we created a new module implementing this interface. As a name for GOBLINT internally we chose `taintPartialContexts` because `taint` was already used, and the name needs to be unique.

Domain

The next step was to choose D and C . According to the concept of our analysis described in Subsection 3.1.1 the domain should be a set of variables. However, we are now analyzing C instead of our toy language. In C not every left-hand side of an assignment is just a simple variable, but can be one of many more complex things e.g. the memory location `*xptr` pointed to by the pointer `xptr`, the fourth place `a[3]` in an array `a`, the member `frac.n` of a struct `frac` and many more. This concept is called Left Value (of an assignment) (`lval`) and there is an implementation of this type provided by GOBLINT in the `Lval.CilLval` module. To be as precise as possible we will use a set of `lvals` instead of a set of variables for the implementation of the taint analysis.

Another point worth mentioning is that we sometimes need the notion of "all variables" (or rather "all `lvals`") when we want to express that everything is tainted. While conceptually using the set X poses no issue, in a concrete implementation this is extremely unpractical and not even realizable if the set is infinitely large. For this case GOBLINT provides a parametrized domain `ToppedSet(Base)`. This domain is either a set of elements of the `Base` type or alternatively a `Top` element which can be interpreted as the "full set of all `Base` elements". Therefore, we finally have $D = \text{ToppedSet}(\text{Lval.CilLval})$ for our domain. Note that this also defines the ordering on the domain to be the regular subset ordering.

It remains to define the module C : We noted in Subsection 3.1.1 that our analysis is inherently context insensitive. Therefore, the context-sensitive part of our analysis C empty, which is expressed with the `Unit` domain provided by GOBLINT.

Transfer functions

These implement the effect of actions on the state, similar to the abstract effects of actions $\llbracket A \rrbracket^\#$ in Chapter 2. Variable declarations are handled by `vdecl` while `branch` handles checks for if-statements and loops. For these two actions our analysis just propagates the state from before, so the two mentioned functions use the default implementation from the `Analysis.IdentitySpec` of GOBLINT.

Much more interesting is the case of the `assign` function which handles the effect of an assignment to an `lval`. For this case we want to add the `lval` to our tainted set. The parameters for the `assign` function are: `ctx` which amongst other things contains the state from before, the `lval` to which a value is assigned and an expression that evaluates to the value that is assigned. We are only interested in `ctx` and the `lval` as to us, only the fact that a value is assigned is relevant and not its value.

Tainting `lvals` is not as straightforward as it might seem at first. Just adding it to the state from before, i.e., the tainted set, only suffices if the `lval` is a specific location in

the memory, e.g., a specific (local or global) variable. The `lval` could however also be a reference to a location in the memory, e.g., a pointer. For these it is not helpful to just taint the reference because we need to know the specific memory locations that are or could be tainted. To solve this issue we make use of GOBLINT's `MayPointTo` query. This takes a reference to the memory and asks all other activated analyses if they have any information about where this reference may point to. Just like everything else in the static analyzer GOBLINT, the answer is an overapproximation, so we can be sure not to miss any location that could be referenced.

In conclusion, tainting an `lval` goes as follows: If the `lval` is a specific memory location, add this `lval` to the tainted set. If it is a reference to the memory described by some expression, send a `MayPointTo` query to ask other analyses which memory locations this expression may point to and add the returned set of `lvals` to the tainted set. We implemented this functionality in a helper function `taint_lval`. Therefore, calling this function is the only thing the `assign` function needs to do as seen in Figure 3.3.

Functions for interprocedural analysis

Function for analysis of multithreaded programs

The query function

We wanted to enable our taint analysis to tell other analyses which `lvals` are tainted at a specific program point. Therefore, we added a new query `MayBeTainted` to the query system of GOBLINT. The result of this query should be a set containing `lvals` which may be tainted, i.e., any `lval` not in the returned set is not tainted.

After this addition we were able to make our `taintPartialContexts` analysis answer this query. Therefore, our analysis implements the query function in such a way that it answers only `MayBeTainted` queries with the current state but does not answer other queries.

3.2.2 Benefiting other analyses

In this section we will discuss how we improved other existing analyses in GOBLINT using the taint analysis we implemented in Subsection 3.2.1. The main analysis that benefited from these changes is the base analysis of the analyzer. This analysis implements a very much extended approach of the basic values-of-variables analysis we formally defined in Chapter 2. The base analysis is however still based on the main goal and basic concept of finding a mapping from program variables to possible values at each program point.




Figure 3.1: Schematic directory structure of GOBLINT. Adapted from [Api14]

- mappings of lvalues new to the caller are taken from callee e.g. newly allocated memory on callee
- mappings not in callee state but are in caller context need to be removed -> in multithreaded programs, if in the caller a mutex was held but unlocked by the callee
- other: keep values from caller if untainted. Lvalues present in the tainted are overwritten with values from callee by folding over tainted set.

- relation analysis (apron) benefited in a similar way
 - mention varEq and condVars for completion??

– Full New Section: ThreadCreate analysis

```
1 module type Spec =
2 sig
3   (* Domain *)
4   module D : Lattice.S
5   module C : Printable.S
6
7   val name : unit -> string
8   val startstate : varinfo -> D.t
9   val query : (D.t, C.t) ctx -> 'a Queries.t -> 'a Queries.result
10
11   (* Transfer functions *)
12   val assign: (D.t, C.t) ctx -> lval -> exp -> D.t
13   val vdecl : (D.t, C.t) ctx -> varinfo -> D.t
14   val branch: (D.t, C.t) ctx -> exp -> bool -> D.t
15
16   (* Functions for interprocedural analysis *)
17   val special : (D.t, C.t) ctx -> lval option -> varinfo -> exp list -> D.t
18   val enter : (D.t, C.t) ctx -> lval option -> fundec -> exp list -> (D.t * D.t) list
19   val return : (D.t, C.t) ctx -> exp option -> fundec -> D.t
20   val combine : (D.t, C.t) ctx -> lval option -> exp -> fundec -> exp list -> C.t option -> D.t
21
22   val context : fundec -> D.t -> C.t
23
24   (* Function for analysis of multithreaded programs *)
25   val threadenter : (D.t, C.t) ctx -> lval option -> varinfo -> exp list -> D.t list
26   val threadspawn : (D.t, C.t) ctx -> lval option -> varinfo -> exp list -> (D.t, C.t) ctx list
27 end
```

Figure 3.2: Simplified Interface for implementing analyses in GOBLINT


```
1 let taint_lval ctx (lval:lval) : D.t =
2   let d = ctx.local in
3   (match lval with
4   | (Var v, offs) -> D.add (v, resolve offs) d
5   | (Mem e, _) -> D.union (ctx.ask (Queries.MayPointTo e)) d
6   )
7
8 let assign ctx (lval:lval) (rval:exp) : D.t =
9   taint_lval ctx lval
```

Figure 3.3: Implementation of the helper `taint_lval` and the `assign` function

4 Evaluation

4.1 Testing

- soundness checked with regression tests from GOBLINT

4.2 Benchmarking

- coreutil as benchmarking programs
- various analysis runs performed with goblint: ctx insensitive with and without taint, precision compared, checks passing compared.

5 Conclusion

Abbreviations

CFG Control flow Graph

lval Left Value (of an assignment)

List of Figures

2.1	Example program (left) and corresponding CFG (right)	2
2.2	Example program (left) and corresponding CFGs for main (middle) and incr (right)	9
3.1	Schematic directory structure of GOBLINT. Adapted from [Api14]	16
3.2	Simplified Interface for implementing analyses in GOBLINT	17
3.3	Implementation of the helper <code>taint_lval</code> and the <code>assign</code> function . . .	18

List of Tables

Bibliography

- [Api14] K. Apinis. “Frameworks for analyzing multi-threaded C.” PhD thesis. Technische Universität München, 2014.
- [ASV12] K. Apinis, H. Seidl, and V. Vojdani. “Side-effecting constraint systems: a swiss army knife for program analysis.” In: *Asian Symposium on Programming Languages and Systems*. Springer. 2012, pp. 157–172.
- [RY20] X. Rival and K. Yi. *Introduction to static analysis: an abstract interpretation perspective*. Mit Press, 2020.