



SCHOOL OF COMPUTATION,
INFORMATION AND TECHNOLOGY —
INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

Combatting the Precision Loss of Partial Contexts in Abstract Interpretation

Felix Sebastian Kraye



SCHOOL OF COMPUTATION,
INFORMATION AND TECHNOLOGY —
INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

**Combatting the Precision Loss of Partial
Contexts in Abstract Interpretation**

**Bekämpfung des Präzisionsverlustes durch
partielle Kontexte in Abstrakter
Interpretation**

Author:	Felix Sebastian Kraye
Supervisor:	Prof. Dr. Helmut Seidl
Advisor:	Michael Schwarz
Submission Date:	15th of February 2023

I confirm that this bachelor's thesis is my own work and I have documented all sources and material used.

Munich, 15th of February 2023

Felix Sebastian Kraye

Abstract

The precision of interprocedural static analyses depends on the variant of context-sensitivity used. While less context-sensitivity grants faster computation times, it comes with a loss in precision. A portion of this precision is actually lost unnecessarily, as some parts of the caller state are not altered during the call, however, they are still overwritten with less precise information from the callee after the call.

We concretize this unnecessary precision loss for values-of-variables analyses and give an approach to recover it. For this, we define a taint analysis tracking which variables may be written by a procedure call. We use this information to update the caller state only with the information about possibly written variables from the callee state after a call and keep the information of definitely unwritten ones. We implement a version of this approach in the Goblint analyzer for the C language and perform benchmarks on it. Additionally, we give a similar approach for a specific thread-related analysis, where the caller state only needs to be updated with the callee state when a thread was created in the procedure call.

The results of our benchmarks show, that actually the precision lost as well as speedup gained through context-insensitivity compared to a fully context-sensitive analysis is rather minuscule for the majority of our benchmark programs. Furthermore, even though our proposed approach recovers a notable portion of the little precision that is lost, it fails to consistently achieve a shorter computation time than a fully context-sensitive analysis. However, we found that the number of timeout and stack overflow errors can be significantly reduced through context-insensitivity. Thus, our approach is best applied in cases, where errors have to be avoided, but precision is more important than computation time.

Contents

Abstract	iii
1 Introduction	1
2 Background	6
2.1 Flow sensitive analysis	7
2.2 Constraint systems	8
2.3 Interprocedural analysis	9
2.4 Context-sensitivity	10
2.5 Partial context-sensitivity	13
2.6 Precision loss	14
3 Combatting the Precision Loss of Variable Analyses	15
3.1 Formal description	15
3.1.1 Taint analysis	15
3.1.2 Improving Variable Analyses	17
3.2 Implementation	18
3.2.1 Taint analysis	19
3.2.2 Benefiting other analyses	26
4 Combatting the Precision Loss of Thread Analyses	32
4.1 Theory	32
4.2 Implementation	36
5 Evaluation	39
5.1 Testing	39
5.2 Benchmarking the improved base analysis	40
5.2.1 SV-Benchmarks	40
5.2.2 GNU Coreutils	43
5.3 Benchmarking the improved threadId analysis	46
6 Conclusions	47
List of Figures	49
Bibliography	50

1 Introduction

Abstract interpretation is a fascinating theory. When its principles are used in static analysis, one can prove certain properties of computer programs. Through the sound abstractions of abstract interpretation, it is ensured that any proven property holds true for all possible executions of a program. The most interesting parts of this application are the different ways by which states of a program can be abstracted and how these abstractions can be combined to gain various kinds of information about a program. The GOBLINT analyzer is a project that applies the principles of abstract interpretation to create a static analyzer [Goba].

This analyzer is specialized in, but not limited to finding concurrency bugs. These are some properties it aims to check:

- Race Detection: Checking that accesses to shared memory never happen simultaneously.
- Assertions and Dead Code: Checking whether specific logical expressions are definitely true at given points within the program.
- Integer Overflows: Verifying that no integer overflows occur in the program.

To gain information about a program, GOBLINT performs various kinds of analyses on the source code. These analyses abstract states of the program in different ways. They are also able to communicate with each other to profit from the information gained by other analyses. Because it is easily expandable, GOBLINT is an interesting framework to try out new approaches in static analysis.

The analyzer is highly configurable. This allows the user to fine-tune the degree of precision they wish. However, it has to be considered that a higher degree of precision usually results in a higher computation time.

With such a configuration option the user can specify a variant of context-sensitivity for each analysis. It is possible to set an analysis to be performed fully context-sensitively, context-insensitively or partially context-sensitively. Context-sensitivity describes the way, in which entry states of functions are differentiated: When full context-sensitivity is chosen, a function is analyzed separately for each entry state, while with context-insensitivity it is only analyzed one time with a summary of all entry states. Different

variants of partial context-sensitivity allow for grouping the entry states in certain ways, where a function is then analyzed once per group with a summary of that group.

```
1  int glob;
2
3  void foo() {
4      //...
5  }
6
7  void main() {
8      glob = 1;
9      foo();
10
11     glob = 2;
12     foo();
13 }
```

Figure 1.1: C code sample with multiple function calls to illustrate the precision loss of partial context-sensitivity.

Consider the program in Figure 1.1. Assume this program is analyzed with the goal to find out which values the program variables can have during program execution. For this, an analysis is used that tracks a set of integers for each variable, i.e., it computes an abstract state describing the possible values of all variables for each program point. The abstract state (in this case a set of values per variable) for a program point is computed by applying the effect of an action, e.g., an instruction, to the state before this action. Most actions have effects that are easily computable, however, function calls are of a more complicated nature, as they can be called from multiple places in a program.

The program in Figure 1.1 contains two calls to `foo()`, one in Line 9 and the other in Line 12. If this program is analyzed context-sensitively, the function is analyzed twice: Once with an abstract entry state describing that "`glob = 1`" and once with a state describing "`glob = 2`". However, if the analysis is performed context-insensitively, both of these two entry states are joined into one abstract state which is used to analyze the function only once. This joined abstract state then has to describe all the concrete states, where "`glob = 1` or `glob = 2`", which is less precise than either of the individual states from before.

When the state after a function call is computed, the information about `glob` is taken from the state returned from the callee. This is because `glob` is a global variable and thus its value can be changed in function calls.

However, consider the case, where `glob` is not changed during the call to `foo()`. Then the information about `glob` in the callee return state is the same as in the entry state for that specific call. This means, that in the case of the context-insensitive analysis, the less precise information "`glob = 1` or `glob = 2`" from the joined entry state is used for `glob` after both calls. This is a loss of precision, considering that the value of `glob` was not altered in the call, and thus, for both states after a call to `foo()` it would be sound to keep the information about `glob` from before that call.

We think this precision loss is avoidable in many cases, where a piece of less precise information is taken from the callee state, even though that piece of information was not changed during the call. Thus, in this thesis, we explore a way to reduce this kind of precision loss of partial contexts in abstract interpretation for certain analyses.

Related work

In the following we discuss some approaches other researchers proposed to improve the precision of (partial) context-sensitivity. The vast majority of these aim to refine context-insensitive analyses with partial or selective context-sensitivity applied to certain elements of a program:

There are multiple approaches proposing variants of "selective context-sensitivity". In general, selective context-sensitive analyses aim to choose a suitable variant of (partial) context-sensitivity depending on the procedure to be analyzed.

A variant of selective context-sensitivity is proposed by Smaragdakis et al. [SKB14]. They introduce an "introspective context-sensitivity" approach for Java programs, where first a context-insensitive analysis is performed. The results of which are used to select program elements that are refined through a context-sensitive analysis of specific methods. In their work, they give different heuristics by which the program elements to be refined are chosen. It is worth mentioning that they acknowledge an often-reported phenomenon: Either context-sensitivity scales rather well in terms of computation time, i.e., it is almost as fast as context-insensitivity, or, in other cases, it scales very badly and takes magnitudes longer. We found a similar phenomenon in our benchmarks.

Li et al.(2018) [Li+18] give a related approach: they also first perform a context-insensitive pre-analysis and use the results to tune context-sensitivity per method. However, instead of choosing between context-sensitivity and -insensitivity, they select from multiple variants of context-sensitivity. These variants are comparable to certain fixed variants of partial context-sensitivity in our terminology.

Another approach is given by Li et al.(2020) [Li+20], that again is concerned with deciding to which methods context-sensitivity should be applied and which variant. However, instead of relying on heuristics like Smaragdakis et al. [SKB14] they identify

program patterns and use these as a basis for the selection. For this, they investigate and identify sources of precision loss caused by context-insensitivity as we do in this thesis. While our focus lies on identifying and combatting the precision loss that occurs with parts of the program state that are not altered during a call, they concentrate on the way precision is lost for parts of the program state that are changed in a call.

Lu and Xue [LX19] give a different approach to selective context-sensitivity: Instead of focussing on choosing whether context-sensitivity or which variant of it should be applied to a method, they tune the context per method on a variable-grained level. For this, they use the concept of "partial context-sensitivity" which we also use in our thesis. However, while we generally fix the variant of partial context-sensitivity for the analysis overall or set it manually, they automatically choose different variants for individual methods. In particular, they use insights from context-free-language reachability at the level of variables to choose which variables should be in the context of a certain method, i.e., with respect to which variables the method should be analyzed context-sensitively. Similar to the publication by Li et al. (2020) [Li+20], they identify and discuss sources of precision loss, but again with the focus on parts of the program state that are changed in a call. The precision loss we address in this thesis occurs in parts of the state that are not changed in a call.

Lastly, Thakur and Nandivada [TN19] propose an interesting approach that makes use of the observation that not all methods have an impact on the caller state. First, the program is pre-analyzed flow-insensitively to identify which parts of the caller state are needed in the callee. The results of this are then used to perform a flow- and context-sensitive analysis on the program, where the analysis of methods that do not impact the caller state is deferred to a post-analysis in order to gain efficiency. In this thesis, we make use of a similar observation, that some parts of the caller state are not impacted by the call. Thakur and Nandivada [TN19] use this insight to compare the contexts for method calls and defer the analysis of certain methods. In contrast to that, we use this observation to identify and keep information from the caller state that is definitely not changed during the call.

Even though the approaches given above are used to analyze the Java programming language, their concepts can be analogously applied to the C language we analyze with GOBLINT in this thesis. However, we mention this approach by Oh et al. [Oh+14], who like us apply their concepts to C programs: They use a context-insensitive pre-analysis to identify which functions are likely to benefit from being analyzed context-sensitively. All these publications propose and discuss different ways to improve context-insensitive analyses in terms of precision. They use various insights and observations to find a suitable variant of context-sensitivity for different program elements or suitable contexts for different methods, functions or procedures. However, the approach we present in this thesis is not concerned with choosing from variants of context-sensitivity or

altering the way contexts are chosen. Instead, it identifies parts of the caller state that are definitely not altered during the call and keeps them from the state before the call. Interestingly, since our concept is concerned with an entirely different aspect than the other approaches we discussed above, it can theoretically be combined with any one of them to improve its precision without interference.

Structure

This thesis is structured as follows: First, we discuss the basics of static analysis in Chapter 2. For this, we introduce constraint systems and how these are used to gain information about the program statically. This is accompanied by the example of a value-of-variables analysis acting on a toy language we use to exemplify applications of static analyses in this thesis. We explain how interprocedural analysis is handled and introduce (partial) context-sensitivity. Here a source of the precision loss is identified. In the following two chapters, we take a closer look at two kinds of analyses that suffer from this loss of precision in different ways. For each kind we propose an approach to reduce the precision loss: In Chapter 3 we aim to improve analyses that track information about variables and in Chapter 4 we give an approach to reduce the precision loss of a thread-related analysis. Both approaches are first discussed conceptually, after which we present the challenges and results of implementing them in the GOBLINT analyzer. To give an evaluation of the proposed approaches, a benchmark of the implementation is performed and inspected in Chapter 5. Our conclusions are presented in Chapter 6.

2 Background

Abstract interpretation is the theory of approximating computer programs in a sound manner. That means that an abstraction might not be absolutely precise, but it definitely is not wrong, i.e., all possible concrete states and properties of a program are described by their abstraction.

An application of abstract interpretation is static analysis. As Rival [RY20] defines it, static analysis is "[...] an automatic technique that approximates in a conservative manner semantic properties of programs before their execution". This means that the program is analyzed just by the given source code without execution. The goal is to prove certain properties about the program in a "sound" manner, i.e., any property that is proven to hold actually does hold. However, from failing to prove a property one cannot conclude that the given property does not hold.

To prove properties, e.g., finding that a program does not contain race conditions or identifying dead code, information about the program has to be gained. This is done by performing various kinds of analyses. We will focus on flow-sensitive analyses in this thesis, i.e., analyses that find properties of the program dependent on the location within it. We will introduce a syntax to formalize flow-sensitive analyses in the following sections. This formalization approach is heavily based on [ASV12].

```
1 void main() {  
2   int x;  
3   x = 0;  
4   if (x == 0)  
5     x = x + 1;  
6 }
```




Figure 2.1: Example program (left) and corresponding CFG (right).

2.1 Flow sensitive analysis

As noted above flow sensitive analyses aim to find properties of the program dependent on the point within the program. Expressed differently this means a flow-sensitive analysis will find an overapproximation of states the program may be in for any given point within the program, from now on called "program point". This state can describe many things dependent on the analysis performed.

First, we define what a program point is: Consider a CFG (Control Flow Graph), where nodes represent points between instructions within the program. Edges are labeled with instructions or checks (from now on collectively called "actions") and describe the transitions between these points (see example Figure 2.1). Then any node on this CFG is what we call a program point.

Concretely let N be the set of all program points. Furthermore, let \mathbb{D} be a domain containing abstract states describing concrete states of the program. This means that some $d \in \mathbb{D}$ can describe many states the program can be in. A concretization function $\gamma : \mathbb{D} \rightarrow 2^{\mathcal{C}}$ can be defined to extract the set of concrete states which are described by an abstract state. Here \mathcal{C} is the set of all possible concrete states.

An analysis is expected to find a mapping $\eta : N \rightarrow \mathbb{D}$ which maps program points to abstract states describing that location within the program, i.e., for $[v] \in N$, $\eta [v]$ should be an abstract state describing all possible states (and possibly more) the program can be in at program point $[v]$.

As an example, we introduce a values-of-variables analysis for integers. This analysis finds a mapping from a set of program variables X to abstractions of their possible values at any given program point. Our toy language supports global variables (globals) as well as local variables (locals). The global variables can be accessed and changed by any procedure, while local ones are only visible to the procedure in which they are declared and can only be accessed and changed by this procedure. Therefore, our set X of variables is the disjoint union of globals G and locals L : $X = G \uplus L$. In the scope of this thesis, we will focus on abstracting integer values by sets of integers. Thereby the goal of our values-of-variables analysis is to find a mapping $X \rightarrow 2^{\mathbb{N}}$ for each program point.

Combining this with the considerations from above, we chose the mapping $\mathbb{D}_v = X \rightarrow 2^{\mathbb{N}}$ as the abstract domain for the values-of-variables analysis. To illustrate what an abstract state from this domain describes, we define the concretization function $\gamma_v : \mathbb{D}_v \rightarrow 2^{\mathcal{C}_v}$. Here a concrete state is a mapping from variables to a single value each: $\mathcal{C}_v = X \rightarrow \mathbb{N}$. Thus, we define the concretization function as follows:

$$\gamma_v M = \{\hat{M} \in \mathcal{C}_v \mid \forall x \in X : (\hat{M} x) \in (M x)\}$$

In summary, the resulting $\eta_v : N \rightarrow \mathbb{D}_v$ for this analysis describes a mapping $\eta_v [v]$ for each program point $[v] \in N$, where $\eta_v [v] x$ is a set containing all values $x \in X$ may possibly hold at $[v]$. From this we can conclude that x cannot hold any value outside $\eta_v [v] x$ at program point $[v]$.

2.2 Constraint systems

We now formulate a way in which we can describe an analysis in the form of constraints. For this we need a partial ordering \sqsubseteq on the domain \mathbb{D} .

Then we create a system of constraints that can be solved for a solution. Consider the edges (u, A, v) of the CFG, where each edge denotes a transition from program point $[u]$ to program point $[v]$ via the action A . Now let each of these edges give rise to a constraint

$$\eta [v] \sqsupseteq \llbracket A \rrbracket^\# (\eta [u])$$

where $\llbracket A \rrbracket^\#$ denotes the abstract effect of the action A defining our analysis. In addition, we need a start state. This is given by $\text{init}^\# : \mathbb{D}$ which is defined depending on the analysis. This gives rise to the start constraint $\eta [s_{\text{main}}] \sqsupseteq \text{init}^\#$ for the starting point of the program $[s_{\text{main}}] \in N$.

We show these ideas with our example of the values-of-variables analysis: For this, a partial ordering \sqsubseteq_v on the domain \mathbb{D} has to be defined. This ordering is needed to formulate the constraints. We define \sqsubseteq_v as follows: A mapping $M_1 \in \mathbb{D}_v$ is ordered below or equal to another mapping M_2 , if and only if for every variable $x \in X$, the set x is mapped to in M_1 is a subset of or equal to the one x is mapped to in M_2 . Formulated formally this is:

$$M_1, M_2 \in \mathbb{D}_v : M_1 \sqsubseteq_v M_2 \iff \forall x \in X : M_1 x \subseteq M_2 x$$

Next, we define the start state $\text{init}^\# = M_\top$ for this domain as the mapping that maps every variable to the full set of integers \mathbb{N} , i.e., $\forall x \in X : M_\top x = \mathbb{N}$. This is because we assume variables (both locals and globals) to be randomly initialized in our toy language.

It remains to define the abstract effect of actions $\llbracket A \rrbracket_v^\#$ for our values-of-variables analysis. For brevity, we just show the effect of a simple variable assignment concretely:

$$\llbracket x = y; \rrbracket_v^\# M = M \oplus \{x \mapsto (M y)\}$$

where $M \oplus \{x \mapsto s\}$ denotes that the mapping M is updated such that x is mapped to the set s in the resulting mapping.

In general, for assignments of expressions e , we evaluate the expression abstractly, i.e., such that the result of the evaluation is the set that contains all integer values e could possibly evaluate to. For example

$$\llbracket x = x + 1; \rrbracket_v^\# \{x \mapsto \{1, 2\}\} = \{x \mapsto \{1, 2\}\} \oplus \{x \mapsto \{2, 3\}\} = \{x \mapsto \{2, 3\}\}$$

A much deeper insight into how expressions are evaluated in abstract interpretation can be found in Patrick Cousot's book "Principles of Abstract Interpretation" [Cou21] in Chapter 3.

2.3 Interprocedural analysis

So far we have only defined how a program without procedure calls is analyzed. Now we want to introduce procedure calls of the form $f()$. For simplicity, we only consider argumentless procedure calls without a return value in our formal descriptions. Arguments and return values can be simulated by using global variables.

Since a call has its own set of local variables to work with and a call stack can contain multiple of the same procedure (e.g. for recursion), procedures are analyzed each in their separate environment. However, we need to consider global variables and how the procedure affects these.

The idea is to give procedures their own starting states and analyze them similarly as we have done before. The final state of the called procedure is then used to be combined back with the state of the caller before the call. Formalized for an edge $(u, f();, v)$ this looks as follows:

$$\begin{aligned} \eta [s_f] &\sqsupseteq \text{enter}^\# (\eta [u]) \\ \eta [v] &\sqsupseteq \text{combine}^\# ((\eta [u]), (\eta [e_f])) \end{aligned}$$

where $[s_f]$ and $[e_f]$ are the start and end nodes of the CFG for procedure $f()$. The functions $\text{combine}^\# : \mathbb{D} \times \mathbb{D} \rightarrow \mathbb{D}$ and $\text{enter}^\# : \mathbb{D} \rightarrow \mathbb{D}$ are defined by the analysis. $\text{enter}^\#$ handles computing the start state for the procedure $f()$, while $\text{combine}^\#$ describes in what way the caller state and the end state of the callee are merged after the call.

It is worth mentioning at this point that even though a procedure can be called from multiple points within the program it is still only analyzed once. For n procedure calls $(u_n, f();, v_n)$ we get n constraints for $[s_f]$: $\eta [s_f] \sqsupseteq \text{enter}^\# (\eta [u_n])$. We can express this differently in a single constraint as follows:

$$\eta [s_f] \sqsupseteq \bigsqcup \{d \mid \exists (u_n, f();, v_n) \in \text{Edges} : \text{enter}^\# (\eta [u_n]) = d\}$$

where \sqcup is the least upper bound, i.e., the least $d \in \mathbb{D}$ according to the ordering \sqsubseteq that is ordered above all of its argument elements. The least upper bound is also called "join".

For our values-of-variables analysis we show how $\text{enter}_v^\#$ and $\text{combine}_v^\#$ are defined. We need to take global variables into account when computing the start state and when combining the caller state with the returned callee state after the call. Therefore, we define the two functions as follows:

$$\text{enter}_v^\# M = M|_{\text{Globals}} \oplus \{x \mapsto \mathbb{N} \mid \forall x \in X\}|_{\text{Locals}_{\text{ce}}}$$

$$\text{combine}_v^\# (M_{\text{cr}}, M_{\text{ce}}) = M_{\text{cr}}|_{\text{Locals}_{\text{cr}}} \oplus M_{\text{ce}}|_{\text{Globals}}$$

where $M|_{\text{Locals}}$ and $M|_{\text{Globals}}$ refers to the mapping M restricted to only the local or global variables respectively. Note that $\text{Locals}_{\text{ce}}$ refers to the locals of the callee while $\text{Locals}_{\text{cr}}$ refers to the locals of the caller.

The $\text{enter}_v^\#$ function first takes the part of the mapping from the caller that contains information about global variables. It adds the local variables used in the procedure to the resulting state. These map to \mathbb{N} because in our toy language they are randomly initialized.

For $\text{combine}_v^\#$ the local part of the caller is kept, but it is updated with the global part of the callee return state. This is done because the latter contains the updated information about global variables after the procedure call.

2.4 Context-sensitivity

In the previous section, we approached the analysis of procedures by analyzing them only once with an abstract start state describing all possible concrete states the procedure could start with. We call this behavior "context-insensitive" as the procedure is analyzed without differentiating between different states with which it is called.

This is not very precise as we will exemplify by applying the values-of-variables analysis to the program in Figure 2.2. We ignore the marked lines of the program for now.

The procedure `incr()` is called twice: Once with $a = 1$ in Line 10 and once with $a = -3$ in Line 13. This leads to two constraints for node $[s_{\text{incr}}]$:

$$\eta_v[s_{\text{incr}}] \sqsupseteq_v \text{enter}_v^\# \eta_v[v_2] = \{a \rightarrow \{1\}\}$$

$$\eta_v[s_{\text{incr}}] \sqsupseteq_v \text{enter}_v^\# \eta_v[v_5] = \{a \rightarrow \{-3\}\}$$

resulting in the joined state $\eta_v[s_{\text{incr}}] = \{a \rightarrow \{-3, 1\}\}$.

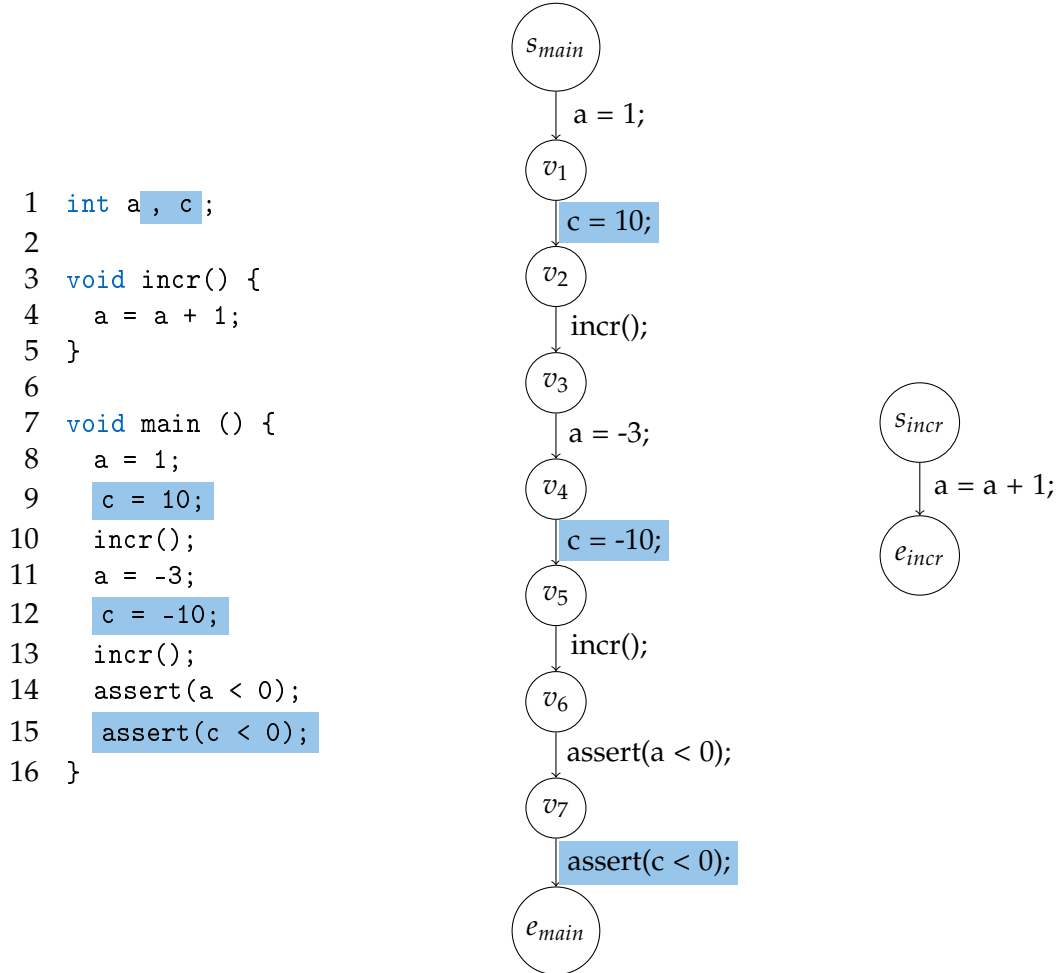


Figure 2.2: Example program (left) and corresponding CFGs for `main` (middle) and `incr` (right). The asserts in Line 14 and Line 15 are proven with a context-sensitive values-of-variables analysis. However, a context-insensitive values-of-variables analysis fails to prove both.

At the end point of the call, the state is $\eta_v [e_{incr}] = \{a \rightarrow \{-2, 2\}\}$, which is then combined with the states of nodes in the main procedure. Therefore, the state at Node $[v_6]$ is $\{a \rightarrow \{-2, 2\}\}$, which is used to check the `assert(a < 0)`; in Line 14. The result of this assertion cannot be determined by the analysis even though it is easy for humans to see that it should hold.

This could be avoided, if the procedure is analyzed twice, once with each entry state. To achieve this, we modify our current approach as follows: Instead of searching for a mapping $\eta : N \rightarrow \mathbb{D}$ we now seek $\eta : (N \times \mathbb{C}) \rightarrow \mathbb{D}$. We call the second part of $(N \times \mathbb{C})$ "context" and \mathbb{C} the "context domain". For now, the context domain is the same as the domain of abstract states, i.e., $\mathbb{C} = \mathbb{D}$.

This allows for different states for the same program point. For now, we differentiate states corresponding to the same program point by the entry state, with which the current procedure was called. Therefore, we adjust the constraints for `enter#` and `combine#` as follows:

$$\begin{aligned} \eta [s_f, \text{enter}^\# (\eta [u, d])] &\sqsupseteq \text{enter}^\# (\eta [u, d]) \\ \eta [v, d] &\sqsupseteq \text{combine}^\# ((\eta [u, d]), (\eta [e_f, \text{enter}^\# (\eta [u, d])])) \end{aligned}$$

With these changes, we can track states for different entry states to procedure calls. In the combine, only the return state that corresponds to the entry state of this specific call is taken into account.

We define the context for the node s_{main} to be M_\top for the initial entry to the program. Thus, the constraint for program initialization is $\eta [s_{main}, M_\top] \sqsupseteq \text{init}^\#$.

We do not have to perform any changes on the values-of-variables analysis to make it context-sensitive. Solely the changes to the general analysis framework above suffice. Applying this changed analysis to the example in Figure 2.2 leads to the procedure `incr()` being analyzed twice with different contexts, assuming we still ignore the marked lines. This leads to the following two entry constraints for different unknowns of the constraint system:

$$\begin{aligned} \eta_v [s_{incr}, \{a \rightarrow \{1\}\}] &\sqsupseteq_v \{a \rightarrow \{1\}\} \\ \eta_v [s_{incr}, \{a \rightarrow \{-3\}\}] &\sqsupseteq_v \{a \rightarrow \{-3\}\} \end{aligned}$$

For node v_6 only the state $\eta_v [e_{incr}, \{a \rightarrow \{-3\}\}] = \{a \rightarrow \{-2\}\}$ is combined with the caller state from before the call. With this information, the analysis can prove that the assertion in the following Line 14 holds for every execution of the program.

2.5 Partial context-sensitivity

While the context-sensitive approach from the previous section might be very precise, it can be quite costly in terms of computation time. To reach a middle ground between a context-insensitive and a fully context-sensitive analysis, we change the approach so that the context is different from the entry state of a call. With this we can group entry states by contexts to analyze a procedure multiple times. This time we group not once per individual entry state, but once per group of entry states.

Thus, we now lift the limitation that $\mathbb{C} = \mathbb{ID}$. This allows for differentiating function calls not by the entry state but by something different defined by the analysis.

To compute the context when entering a procedure, we define a new function $\text{context}^\# : \mathbb{ID} \rightarrow \mathbb{C}$. Additionally, the constraints for $\text{enter}^\#$ and $\text{combine}^\#$ are changed as follows:

$$\eta [s_f, \text{context}^\# (\eta [u, c])] \sqsupseteq \text{enter}^\# (\eta [u, c])$$

$$\eta [v, c] \sqsupseteq \text{combine}^\# ((\eta [u, c]), (\eta [e_f, \text{context}^\# (\eta [u, c])]))$$

for an edge $(u, f();, v)$.

This formalization results in multiple constraints for a single contextualized starting variable $[s_f, c']$. We can alternatively formulate this as

$$\eta [s_f, c'] \sqsupseteq \bigsqcup \{ \text{enter}^\# (\eta [u_n, c_n]) \mid \exists (u_n, f();, v_n) \in \text{Edges} : \text{context}^\# (\eta [u_n, c_n]) = c' \}$$

i.e., the constraint for the variable $[s_f, c']$ is the least upper bound of all entry states for some call of f , which have the same context c' as the constraint variable. Or expressed differently, all states computed by $\text{enter}^\# d$ for f are grouped by the context $c' = \text{context}^\# d$, where each group is joined by \bigsqcup to produce a constraint for a starting variable $[s_f, c']$ with the respective context.

With this formal model, we have the option to perform an analysis completely context-sensitively ($\mathbb{C} = \mathbb{ID}$ and $\text{context}^\# = \text{enter}^\#$), completely context-insensitively ($\mathbb{C} = \{\bullet\}$) or anything in between. Note that we define $\{\bullet\}$ as the "unit domain" which contains exactly one element with the trivial ordering $\bullet \sqsubseteq \bullet$.

We have to note here that there are some severe issues with the approach for (partially) context-sensitive analyses described in this thesis: The resulting system of constraints may not be finite and some variables in the constraint system may depend on an infinite number of other variables. Thus, it is computationally complex to compute a solution to the system of constraints. However, we stick with the described methodology in the scope of this thesis for simplicity. An extension to the approach which addresses the mentioned issue can be found in [ASV12].

2.6 Precision loss

The main source of the precision loss in context-insensitive or partially context-sensitive analyses is the join over all states with the same context, i.e., when we take the least upper bound of a group of entry states. Consider a procedure f that has no effect, i.e., $s_f = e_f$. Even for this procedure, the $\text{combine}^\#$ function receives the less precise result of the join \sqcup to combine it with the caller state. In this simple case, the result would be more precise if the $\text{combine}^\#$ function could directly use the result from the corresponding $\text{enter}^\#$ as the callee state for combining.

Even for procedures that do change the state, there might be some parts of the state which are untouched by the call. If we can identify these untouched parts, we could reduce the precision loss experienced by using partial contexts.

We clarify the source of the precision loss mentioned above with an example: For this, we once again consider the example program Figure 2.2. This time we take the marked lines into account. When the program is analyzed context-insensitively with the values-of-variables analysis, not only does the state at the start node for $\text{incr}()$ s_{incr} represent two possible values for the global variable a , but also for c . Therefore, the state for this node is

$$\eta_v [s_{\text{incr}}, \bullet] = \{a \rightarrow \{-3, 1\}, c \rightarrow \{-10, 10\}\}$$

Even though the variable c is never changed within $\text{incr}()$, the mapping $c \rightarrow \{-10, 10\}$ is still copied into the caller state when combining the states for node v_6 . Thus, the information gained by the context-insensitive values-of-variables analysis does not suffice to determine the assertion in Line 15 to hold. This loss of precision could easily be avoided if we had some idea of which global variables are definitely not changed by a procedure call.

3 Combatting the Precision Loss of Variable Analyses

In this chapter, we describe an approach to reduce the precision loss described in Section 2.6 for analyses tracking information about variables. We call analyses of this kind "variable analyses".

First, we use the syntax for flow-sensitive analyses from Chapter 2 to formally explain the idea. After that, we explain a concrete implementation of the approach into the GOBLINT analyzer.

3.1 Formal description

We want to reduce the mentioned precision loss for variable analyses. The basic idea to achieve this is to track for each procedure which variables have been written or have possibly been altered in some other way within that procedure. This information is then used in variable analyses when combining the abstract state from the caller with the abstract return state given by the callee at the end of the procedure. We will exemplify this with the values-of-variables we introduced in Chapter 2.

In the following, we call a variable that has been written or altered in the current procedure context "tainted". Therefore, we introduce a new taint analysis tracking which variables have been tainted within the context of the current procedure in the following section. It is worth mentioning that our notion of taintedness is related but different from other uses of the term "taint analysis".

3.1.1 Taint analysis

In this section we define the syntax for the taint analysis we propose in this thesis: Since we want to find a collection of tainted variables per program point, a suitable domain for this analysis is the powerset of the set of variables X ordered by the subset relation:

$$\mathbb{D}_t = 2^X \text{ with } \sqsubseteq_t = \subseteq$$

To be compatible with the notion of partially context-sensitive analyses from Section 2.5 we also need to specify a context domain \mathbb{C}_t which we define later. Note that we seek to compute a mapping from program points (with context) to sets of variables, i.e., $\eta_t : (N \times \mathbb{C}_t) \rightarrow \mathbb{D}_t$. To interpret this with the goal of our taint analysis in mind, we note that $\eta_t [n, \bullet] = T$ denotes that T is the set of possibly tainted variables at program point n . Expressed differently this means that for any variable $x \in T$ we cannot exclude that this variable was altered from the start of the current procedure up until the program point n . Note here that the tainted set T not only includes variables that have been tainted by statements of the current procedure, but also variables that have been tainted within other procedures called by the current one.

It remains to define \mathbb{C}_t , $\text{init}_t^\#$, $\text{enter}_t^\#$ and $\text{combine}_t^\#$ as well as the abstract effects of actions $\llbracket A \rrbracket^\#$. Recall that the notion of a "tainted" variable is defined in relation to the current procedure. This means we want to start without any variable being initially tainted when entering a procedure. It is worth pointing out that the entry to a procedure call does not depend on the state where it is called. Therefore, we design our analysis to be context-insensitive, i.e.,

$$\mathbb{C}_t = \{\bullet\} \text{ and trivially } \text{context}_t^\# T = \bullet$$

With these considerations we can also define $\text{enter}_t^\#$ and $\text{init}_t^\#$ as follows:

$$\text{enter}_t^\# T = \text{init}_t^\# = \emptyset$$

It is worth pointing out here that the function $\text{enter}_t^\# T$ is always equal to the empty set regardless of its argument T . Therefore, it computes the same entry state for each call of a certain procedure.

When combining the caller state with the returned callee state, we note that we need to keep the tainted set from before the call, as a tainted variable can never get "untainted" again, no matter what the procedure does. Additionally, we add the tainted set returned by the callee, since anything tainted in the call needs to be considered tainted after the call as well. This is because we want to know which variables have been altered in a procedure call, no matter if the tainting happened within the procedure itself or within a further procedure call. This leaves us with the following equation for the $\text{combine}_t^\#$ function:

$$\text{combine}_t^\# (T_{cr}, T_{ce}) = T_{cr} \cup (T_{ce} \setminus \text{Locals}_{ce})$$

Note that we removed the callee local variables Locals_{ce} because these are not accessible by the caller and all of its callers anyway, so it is not useful to keep track of them.

Lastly, we define the abstract effects of actions. Most of these (including checks) do not do anything besides propagating the state from before. The only major exception

are variable assignments. For these, we note that the specific variable, to which the value is assigned is added to the tainted set. This is independent of the expression that evaluates to the assigned value, as we are only interested in the fact that the variable on the left of the assignment is altered. This leaves us with the following abstract effects of actions:

$$\llbracket A \rrbracket^\# T = \begin{cases} T \cup \{x\} & \text{if } A \equiv (x = e;) \\ T & \text{else} \end{cases}$$

where e is any arbitrary expression.

This concludes our definition of the taint analysis.

3.1.2 Improving Variable Analyses

In this section, we see how the information from the taint analysis helps us to improve variable analyses. We show this with the example of the values-of-variables analysis we introduced in Chapter 2.

Recall the source of the precision loss we want to reduce as described in Section 2.6. This loss occurred when a global variable was updated with a less precise value after a procedure call even though this specific variable was not changed by the call.

Thanks to the taint analysis we defined in the previous section, we now have a way to get information about which variables can be altered by a procedure $f()$ and which surely stay untouched. The latter of which are exactly those variables that are not in the tainted set of the end node e_f for that procedure.

With this insight we can now update the $\text{combine}_v^\#$ function of our values-of-variables analysis as follows:

$$\text{combine}_v^\#(M_{cr}, M_{ce}) = M_{cr}|_{Locals_{cr} \cup (Globals \setminus T_{ce})} \oplus M_{ce}|_{Globals \cap T_{ce}}$$

where $T_{ce} = \eta_t[e_f, c]$ for an edge $(u, f();, v)$ and the corresponding context c .

Similar to before, the $\text{combine}_v^\#$ function takes the caller mapping, restricts it to a subset of caller reachable variables, and updates this mapping with the callee mapping restricted to the rest of caller reachable variables. In other words, the caller reachable variables are partitioned into two sets such that one subset is taken from the caller state while the other one is taken from the callee state. Before our change, the partitioning was done strictly in such a way that only the caller local variables were taken from the caller state and all global variables from the callee state. With our improvement, the global variables that are not tainted by the callee are also taken from the caller state and not from the callee anymore. Thereby the precision loss for untainted variables is eliminated.

One might wonder if our change could lead to a case, where the callee state has a more

precise value for a variable that is discarded because this variable is not in the tainted set. Concretely this situation would be described by

$$\exists \text{ Edge } (u, f();, v), x \in \text{Globals} : x \notin \eta_t [e_f, \bullet] \wedge (\eta_v [e_f, \bullet] x \subsetneq \eta_v [u, \bullet] x)$$

where we use the subset orderings \subseteq and \subsetneq to compare the sets of integer values, x could hold at different program points.

From $x \notin \eta_t [e_f, \bullet]$ we know that x has not been altered in the procedure $f()$ since the node s_f up until e_f , and therefore it holds that

$$\eta_v [e_f, \bullet] x = \eta_v [s_f, \bullet] x$$

Because x is a global we get with the definitions of \sqsubseteq_v and $\text{enter}_v^\#$:

$$\eta_v [s_f, \bullet] x \supseteq (\text{enter}_v^\# (\eta_v [u, \bullet])) x = \eta_v [u, \bullet] x$$

Therefore, $\eta_v [e_f, \bullet] x \supseteq \eta_v [u, \bullet] x$ which is a contradiction to the proposed case which we can therefore exclude.

3.2 Implementation

Before explaining the process of implementing the proposed taint analysis and its usage to improve other analyses, we introduce the GOBLINT analyzer and its structure in this paragraph. The core functionality of GOBLINT is to statically analyze C programs using an approach similar to the one described in Chapter 2. This generally works as follows: The C input file is preprocessed with `goblint-cil`, an adaption of the "C Intermediate Language (CIL)" project. CIL is a front-end that parses and compiles C programs into a simplified subset of C [Cil]. For example, it moves assignments that happen within conditional checks out of the check and puts them before it. A CFG is generated from the output of the CIL front-end. This graph is then used together with the specifications of various analyses to generate a constraint system. GOBLINT solves this constraint system and produces different kinds of outputs to the user according to the solution (e.g. notifications, warnings or a visualization of the full solution).

It is worth mentioning that GOBLINT can perform multiple analyses on a program at the same time. For this, a compound domain is built (for the domain of states as well as for the context domain), which is a tuple of the domains corresponding to the analyses to be performed. To generate constraints, all activated analyses are taken into account, where the specification of each analysis acts on its corresponding part of the compound domain. Information can be transferred between the different analyses via a system




Figure 3.1: Schematic directory structure of GOBLINT. Adapted from [Api14].

called "queries".

Figure 3.1 shows the inner structure of the analyzer. It shows, that GOBLINT provides parametrized domains that can be used in the specifications of the analyses, e.g., Set and Map. It is also shown that multiple analyses are then compounded into one Master Control Program (MCP) that is used by the framework to generate constraints from the CFG. The resulting constraint system is then solved by one of the solvers GOBLINT provides in /solvers.

For a deeper insight into the inner workings of GOBLINT refer to [Api14].

3.2.1 Taint analysis

To define an analysis the GOBLINT analyzer provides an interface, of which the most relevant parts can be seen in Figure 3.2. This interface requires two modules *D* and *C* which correspond to the domain of abstract states and the context domain. It also requires the following functions to be defined by an analysis:

- `name` to uniquely refer to an analysis.
- `startstate` to define the state used when entering the analysis (similar to `init#` from Chapter 2).

- query to implement the query system of GOBLINT. This allows an analysis to broadcast information to be used in analyses.
- *Transfer functions* which define the abstract effects of actions (similar to $\llbracket A \rrbracket^\#$ from Chapter 2).
- *Functions for interprocedural analysis*
- *Functions for the analysis of multithreaded programs*

For our taint analysis we create a new module implementing this interface.

As a name for GOBLINT internally we chose `taintPartialContexts` because `taint` was already used, and the name needs to be unique. In the following we explain in detail how our new analysis module implements the given interface:

Domain

We need to choose the domain of abstract states D and the context domain C . According to the concept of our analysis described in Section 3.1.1, the domain should be a set of variables. However, we are now analyzing the C language instead of our toy language. In C not every left-hand side of an assignment is just a simple variable, but can be one of many more complex things, e.g., the memory location `*xptr` pointed to by the pointer `xptr`, the fourth place `a[3]` in an array `a`, the member `frac.n` of a struct `frac` and many more. All these options are described by a concept called "lval" (Left Value (of an assignment)). There is an implementation of this type provided by GOBLINT in the `Lval.CilLval` module. To be as precise as possible we use a set of lvals instead of a set of variables for the implementation of the taint analysis.

Another point worth mentioning is that we sometimes need the notion of "all variables" (or rather "all lvals") when we want to express that everything is tainted. While conceptually using the full set X poses no issue, in a concrete implementation this is extremely impractical and not even realizable if the set is infinitely large. For this case, GOBLINT provides a parametrized domain `ToppedSet(Base)`. This domain is either a set of elements of the `Base` type or a `Top` element which can be interpreted as the "full set of all `Base` elements". Therefore, we finally have $D = \text{ToppedSet}(\text{Lval.CilLval})$ for our domain. Note that this also defines the ordering on the domain to be the regular subset ordering.

It remains to define the module C : We noted in Section 3.1.1 that our analysis by itself is context-insensitive. Therefore, the context domain C of our taint analysis is empty, which is expressed with the `Unit` domain provided by GOBLINT. Note here, that this does not mean that the taint analysis is always performed context-insensitively, i.e., it is only performed once per function. Since GOBLINT uses a compound domain,

```
1 module type Spec =
2 sig
3   (* Domain *)
4   module D : Lattice.S
5   module C : Printable.S
6
7   val name : unit -> string
8   val startstate : varinfo -> D.t
9   val query : (D.t, C.t) ctx -> 'a Queries.t -> 'a Queries.result
10
11   (* Transfer functions *)
12   val assign: (D.t, C.t) ctx -> lval -> exp -> D.t
13   val branch: (D.t, C.t) ctx -> exp -> bool -> D.t
14
15   (* Functions for interprocedural analysis *)
16   val special : (D.t, C.t) ctx -> lval option -> varinfo -> exp list ->
       D.t
17   val enter : (D.t, C.t) ctx -> lval option -> fundec -> exp list ->
       (D.t*D.t) list
18   val return : (D.t, C.t) ctx -> exp option -> fundec -> D.t
19   val combine : (D.t, C.t) ctx -> lval option -> exp -> fundec -> exp
       list -> C.t option -> D.t -> Queries.ask -> D.t
20
21   val context : fundec -> D.t -> C.t
22
23   (* Functions for the analysis of multithreaded programs *)
24   val threadenter : (D.t, C.t) ctx -> lval option -> varinfo -> exp list
       -> D.t list
25   val threadspawn : (D.t, C.t) ctx -> lval option -> varinfo -> exp list
       -> (D.t, C.t) ctx -> D.t
26 end
```

Figure 3.2: Simplified Interface for implementing analyses in GOBLINT

there might be other context-sensitive analyses, forcing the whole compound analysis to analyze a function multiple times. Our taint analysis however never contributes differing sub-contexts to the compound context.

The `startstate` function

This function computes the initial state for our analysis similar to the `init#` function we introduced in Chapter 2. As discussed in Section 3.1.1, we implement this function so that it returns the empty set.

We note here, however, that in practice we do not use the way, `startstate` is implemented in the scope of our thesis. In GOBLINT, this function is called before the main function is even entered. Thus, `enter` (which we define later) is still used to compute the entry state for the main function. We chose to implement `startstate` in this way for consistency.

Transfer functions

These functions implement the effect of actions on the state, similar to the abstract effects of actions $\llbracket A \rrbracket^\#$ in Chapter 2. For example, `branch` handles checks for if-statements and loops. For this and most other actions, our analysis just propagates the state from before, so they use the default implementation from the `Analysis.IdentitySpec` of GOBLINT. The default implementations of `Analysis.IdentitySpec` propagate the given state without change for any action.

Much more interesting is the case of the `assign` function which handles the effect of an assignment to an `lval`. For this case, we want that the `lval` is added to the tainted set. The parameters for the `assign` function are `ctx` which amongst other things contains the state from before, the `lval` to which a value is assigned and an expression that evaluates to this value that is assigned. We are only interested in `ctx` and the `lval`, as for the taint analysis only the fact that a value is assigned is relevant and not its concrete value.

Tainting `lvals` is not as straightforward as it might seem at first. Just adding it to the state from before, i.e., the tainted set, only suffices if the `lval` is a specific location in the memory, e.g., a specific (local or global) variable. The `lval` could however also be a reference to a location in the memory, e.g., a pointer. For these, it is not helpful to just taint the reference because we need to know the specific memory locations that could be tainted. To solve this issue, the analysis makes use of GOBLINT's `MayPointTo` query. This takes a reference to the memory and asks all other activated analyses if they have any information about where this reference may point to. Just like everything else in the static analyzer GOBLINT, the answer is an overapproximation, so we can be sure not to miss any location that could be referenced.

```

1 let taint_lval ctx (lval:lval) : D.t =
2   let d = ctx.local in
3   (match lval with
4   | (Var v, offs) -> D.add (v, resolve offs) d
5   | (Mem e, _) -> D.union (ctx.ask (Queries.MayPointTo e)) d
6   )
7
8 let assign ctx (lval:lval) (rval:exp) : D.t =
9   taint_lval ctx lval

```

Figure 3.3: Implementation of the helper `taint_lval` and the `assign` function for the taint analysis in GOBLINT.

In conclusion, tainting an `lval` goes as follows: If the `lval` is a specific memory location, this `lval` is added to the tainted set. If it is a reference to the memory described by some expression, a `MayPointTo` query is sent to ask other analyses which memory locations this expression may point to and add the returned set of `lvals` to the tainted set. We implemented this functionality in a helper function `taint_lval`. Therefore, calling this function is the only thing the `assign` function needs to do as seen in Figure 3.3.

Functions for interprocedural analysis

Here we define the functions `context`, `enter` and `combine`. These functions work similarly to their abstract counterparts as described in Chapter 2. In addition to these known functions, the interface also requires two additional functions: `return` which handles return statements right before a function is left and `special` which handles calls to library functions or other functions, for which the source code is not available to be analyzed.

Our implementation does not differ a lot from the proposed formal description in Section 3.1.1. Since we are analyzing C and not our toy language, the only major difference is that return values and function arguments have to be handled. Therefore, we implement these functions as follows:

`context:`

Our context domain is the `Unit` domain, and we do not want to generate different contexts for this analysis. Thus, this function always returns the unit element.

enter:

Like we discussed in Section 3.1.1, this function always returns the empty set, which is the entry state for the called function.

combine:

The combine function first checks if there is an lval to which the return value is assigned. If so, it taints this respective lval in the caller state using the helper function `taint_lval` introduced in the "Transfer Functions" section above. After that, it computes the union of the resulting state with the returned callee state and returns it.

Summarized, the result of this function is the union of both states it receives for combining. It additionally adds lvals which are possibly tainted through the assignment of the return value.

return:

In our formal description in Section 3.1.1, the `combinet` function removed variables unreachable by the caller. In the concrete implementation, we give this functionality to the return function, so the removal happens right before the combine. Function arguments are also removed, as these are unreachable by the caller similar to local variables. It is worth pointing out that not just all lvals corresponding to local variables or arguments are removed. A function might exist multiple times in the current call stack, e.g., when the function is recursive. This can result in the existence of multiple versions of the same local variable. GOBLINT treats these as being the same variable. Thus, we risk that, when a local variable is removed, this also removes a different version of the variable lower in the call stack, for which the taintedness information is still needed. To address this issue, return sends an `IsMultiple` query for each variable to be removed and only removes those, that surely do not have multiple versions. This query is already provided by GOBLINT.

special:

This function addresses library functions or other functions, for which the source code is not available to be analyzed. The simple way to handle these is to just return Top, i.e., saying "everything could be tainted", after a special call.

This is how unknown functions are handled, however, GOBLINT provides "Library Descriptors", which contain information about some known C library functions, e.g., `printf`, `malloc`, `cos`, etc. With the respective Library Descriptor of a function, information can be gained about which addresses are "shallowly" written and which are "deeply" written by the call. Shallowly written addresses point to lvals which might be directly written. Deeply written addresses however point to lvals where not only the

lval itself but anything it might recursively point to could be written. Therefore, the special function makes use of GOBLINT's `MayPointTo` and `ReachableFrom` queries in the following way:

First, the function checks if a Library Descriptor is available. If not, `Top` is returned. Otherwise, the shallowly and deeply written addresses are obtained from the Descriptor. Consequently, the union of

- the state before the call
- anything that is possibly tainted by the return value (using `taint_lval` like in `combine`)
- the set of lvals returned by the `MayPointTo` query for any shallowly written address
- the set of lvals returned by the `ReachableFrom` query for any deeply written address

is returned by the special function.

Functions for the analysis of multithreaded programs

To be able to analyze multithreaded programs, GOBLINT's analysis interface requires the following functions: `threadenter` to compute the starting state for newly created threads and `threadspawn` which computes the effect of a thread creating instruction to the state of the creating thread.

We implement the former of these two functions similarly to our `startstate` and `enter`. Thus, `threadenter` returns the empty set.

To implement the other function, `threadspawn`, we consider how a thread creation affects the state of the creator. We note that the only relevant effect for our notion of taintedness is, that the thread creating function may write thread ID variables to which it receives a reference as an argument. Thus, this function uses the helper function `taint_lval` defined in the "Transfer Functions" section to add possibly tainted lvals to the state from before and returns the result.

The query function

We want to enable our taint analysis to tell other analyses which lvals are tainted at a specific program point. Therefore, we add a new query `MaybeTainted` to the query system of GOBLINT. The result of this query should be a set containing lvals which may be tainted, i.e., any lval that is not in the returned set is definitely untainted.

After this addition, we make our `taintPartialContexts` analysis answer to this query. Therefore, this analysis implements the query function in such a way that it answers only to `MaybeTainted` queries with the current state but does not answer other queries.

3.2.2 Benefiting other analyses

In this section, we discuss how we improved other existing analyses in GOBLINT using the taint analysis we implemented in Section 3.2.1.

Improving the `base` analysis

The main analysis benefitting from the taint analysis is the `base` analysis of GOBLINT. This analysis implements a very much extended approach of the basic values-of-variables analysis we formally defined in Chapter 2. The `base` analysis is however still based on the main goal and basic concept of finding a mapping from program variables to possible values at each program point. Therefore, this analysis uses a mapping from variables to their possible values as part of its domain. However, here the `ValueDomain` of the mapping is much more complex than just a set of possible integers. It provides abstractions for virtually any type in C, including arrays, structs and pointers. Even more though, the `ValueDomain` is highly configurable. Amongst other options, it allows choosing between different ways of abstracting integer values or arrays. One interesting option related to the topic of this thesis is the possibility to choose between different variants of context-sensitivity: the analysis can be fully context-sensitive, insensitive with respect to integer variables (abstracted by intervals or in general), only sensitive with respect to pointers or completely insensitive. When choosing anything but the completely context-sensitive option, this analysis experiences the (avoidable) loss of precision described in Section 2.6.

To reduce this loss we need to change the `combine` function of the `base` analysis so that it uses the results of our `taintPartialContexts` analysis. To get a better understanding of what needs to be changed, we first describe how the `combine` function was implemented before our changes:

1. The return value is saved. Its value is removed from the callee state.
2. All globals are removed from the caller state.
3. Everything from the callee state is added to the caller, possibly overwriting caller values. This excludes the return value which is handled separately.
4. Some further adjustments are performed to the resulting state according to the configuration.

5. The saved return value is added to the state before the state is returned.

To implement our changes we will focus on steps 2 and 3, where the caller mapping is updated. The other steps will remain the same.

The core idea to implement the concept proposed in Section 3.1.2 is as follows: First, the set of possibly tainted lvals is obtained from the callee. Then the algorithm iterates over its elements one by one, where for each tainted lval the caller mapping is updated with the corresponding value from the callee mapping, i.e., the algorithm gets the value corresponding to that lval from the callee mapping and sets the lval to map to this value in the caller mapping. This functionality of updating the caller mapping with the callee mapping using the tainted set is implemented in a helper function `combine_st`.

Before explaining how this helper function is implemented, we first show how it is embedded in the current implementation of the `combine` function. As discussed, we alter steps 2 and 3: First a `MayBeTainted` query is sent to the return state of the callee. We then check if the query returned the Top set, i.e., the notion that everything is tainted. In this case, we perform steps 2 and 3 just like before without change. Amongst other cases, this can happen, when the base analysis is run without our taint analysis being activated.

We now define what happens, when the result of the query is an explicit set. Before calling the helper function `combine_st`, we have to handle two special cases here:

- For a global variable, there is no mapping in the callee state, but there is one in the caller state. This case can occur in multithreaded mode, if this variable was protected by a mutex before the call, but the mutex was released in the called function. In this case, the mapping for this variable would be removed from the state within the callee. In the `combine` function, we do need to keep this new information from the callee for such a variable, i.e., remove it from the caller mapping. Therefore, we filter over the caller mapping and remove all globals, for which there exists no mapping in the callee mapping.
- For a global variable, there is a mapping in the callee state, but there is none in the caller state. This case can occur if new information is gained within the call, e.g., some new memory is allocated. This information is not tracked by the tainted set and would therefore not be copied into the caller state. Since we still want to have this new information after the combine, we add all these mappings from the callee to the caller.

These cases have to be handled separately, as for these the corresponding lval is not necessarily contained in the tainted set. After the two special cases are handled, we use the `combine_st` helper function to finally update the tainted lvals in the caller state. We then proceed with the resulting state to steps 4 and 5 like before.

We note here that we added a new parameter `f_ask` to the `combine` function. To do this we had to update the analysis interface and consequently all analyses implementing it. This new parameter allows the analyses to send queries to the returned callee state, which was not possible before.

`combine_st:`

This helper function takes the caller state (updated according to the two special cases), the callee state and the set of tainted lvals. The difficulty here is, that while the tainted set is a set of lvals, the mappings from the states of the base analysis are mappings from variables to abstractions of their possible values. This means, that our tainted set may include specific places in an array or specific members of a struct, e.g., `a[3]` and `frac.n`. In contrast to this, the mappings we want to combine do not map `a[3]` and `frac.n` to abstract values, but rather map `a` to some abstraction of an array and `frac` to some abstraction of a struct. To solve this issue we make use of the `get` and `set_savetop` functions provided by the module of the base analysis. With these functions, it is possible to get and set values of addresses to specific lvals in a variable mapping of the base analysis.

Therefore, the implementation of this function goes as follows: We fold over the tainted set. For each lval we build an address to this lval. Then we try to get the value this address points to from the callee state. If this returns a value, we use `set_savetop` to update the caller state, i.e., set the address to the value we got. Otherwise, we proceed with the next lval.

There are however a few special cases to handle: One issue is that in GOBLINT there exists a domain for abstracting array values called "partitioned array". This abstraction saves an index that it uses to split an array into three parts: The group of all values to the left of the index, the value at the index itself and the group of all values to the right. Each of the three parts is abstracted with a collective value. The index can be either a specific integer or a variable.

For array variables abstracted with this "partitioned array" domain, copying lvals one by one does not work, as the information of the partitioning is lost when it is attempted. Therefore, we check if the current lval to update corresponds to a place in an array abstracted with the "partitioned array" domain and if so, we copy the whole partitioned array from the callee mapping to the caller mapping.

A similar issue occurs with values of void type, as for these the `get` function does not work correctly. Therefore, we get the value for the corresponding variable from the callee mapping itself and update the caller mapping with it.

The last issue we need to address is again related to partitioned arrays. Recall, that an array can be partitioned by the value of a variable. This means, that if a variable

is tainted which is used as a partitioning index, the partitioned array in the caller mapping is invalid. Therefore, for each `lval` we check if it corresponds to a variable partitioning an array. If so, we update the caller state by copying all abstracted arrays which are partitioned by the variable in question from the callee mapping to the caller mapping. This is possible because the state of the base analysis keeps track of which arrays are partitioned by certain variables.

Finally, after we are done folding over all tainted `lvals`, the function returns the modified caller state.

Improving other Variable Analyses

In GOBLINT there are multiple other analyses that can be described by our notion of a variable analysis. For some of these, we can reduce the precision loss of partial contexts by using the results from our `taintPartialContexts` analysis. We briefly discuss the details of these improvements in this section:

The `varEq` analysis:

This analysis tracks, which `lvals` definitely hold the same value, regardless of what this value is. As an example, after an assignment `a[3] = y`; this analysis propagates the information that `a[3]` and `y` in an equality set, until either `a[3]` or `y` are written.

One can construct a case, where a function `f()` is called twice, once where some equality `x = y` holds and once where it does not hold. We assume that `x` and `y` are not altered within `f()`. If the `varEq` analysis is performed context-insensitively, we would lose the equality information when the `combine` is performed after both calls. This is because the equality information is discarded when the entry states are joined.

So far, the `combine` function just propagated the callee return state. This can be improved in the following way: We obtain the tainted set and remove all equality sets containing at least one tainted `lval` from the caller state. Then we compute the greatest lower bound of the resulting caller state and the callee state. This means we compute a state that unifies the information from the callee state with the information from the caller state minus tainted `lvals`.

In the example from above this changed `combine` allows the analysis to keep the equality `x=y` when the taint analysis is activated.

The `relation` analysis:

This analysis tracks relations between variables. The analysis can be configured to use different kinds of relations implemented in different domains. An example of these domains is the octagon domain, which works with relations of the form $\langle x \rangle + \langle y \rangle \geq 0$, where $\langle x \rangle$ stands for either x or $-x$. The relation analysis is parametrized in such a way

that our changes apply to the analysis independent of the selected relation domain. The case where this analysis unnecessarily loses precision because of context-insensitivity is similar to the one discussed in the paragraph discussing the improvement of the `varEq` analysis: A function `f()` is called twice, where some relation between the variables `x` and `y` once holds and once does not hold. When the entry states for both calls are joined, the relation in question is removed and is therefore missing after the call. We once again assume that `x` and `y` are not altered within `f()`.

Before we apply our changes, the `combine` function of the `relation` analysis is implemented so that it removes all information related to global variables from the caller state and merges the result with the callee state. The result contains the relations from the caller that are not related to globals and all relations from the callee.

By accessing the result of our taint analysis via a query, we can improve this way of combining. We achieve this by only removing relations related to tainted global variables and keeping those that only relate to locals and untainted globals. After that, the result is merged with the callee state like before.

We note here, that this analysis tracks relations between variables while our tainted set contains tainted `lvals`. Therefore, we convert the set of tainted `lvals` to a set of tainted variables. This is easily possible since all `lvals` in our tainted set refer to variables or specific parts of variables. In particular, they do not point to some memory location, which would make identifying a corresponding variable difficult.

The `condVars` analysis:

This analysis tracks equalities between variables and logical expressions. Take the following statement as an example: `tv = (c == 0);`. For this statement, the `condVars` analysis tracks, that the variable `tv` holds the value of the logical expression `c == 0`. This information can be used, when there is an `if` statement, checking whether `tv` is true or not. Knowing `tv` is equal to `c == 0`, this analysis can provide the information that `c == 0` evaluates to `true` in the `true`-branch of the `if` statement.

Currently, the `combine` function of this analysis is implemented to discard the callee return state and remove all information related to global variables from the caller state. This results in the loss of all information related to globals whenever a function is called, regardless of whether the `condVars` analysis is performed context-sensitively or -insensitively.

We can improve this by only removing information related to tainted globals from the caller state and keeping information related to untainted globals. This makes the `condVars` analysis more precise whenever the `taintPartialContexts` analysis is activated. Again we note here, that this analysis is concerned with information related to variables, not `lvals`. Therefore, we convert the set of tainted `lvals` to a set of tainted

variables like in the previous paragraph.

4 Combatting the Precision Loss of Thread Analyses

To track relational information between variables in multithreaded programs, M. Schwarz et al. propose a set of analyses specialized for this purpose in their paper "Clustered Relational Thread-Modular Abstract Interpretation with Local Traces" [Sch+23]. In this chapter, we take a closer look at one of these analyses and identify a potential loss of precision when this analysis is performed partially context-sensitively. Furthermore, we propose an approach to reduce this precision loss in certain cases.

4.1 Theory

In Section 6 of their paper [Sch+23] they propose an analysis that identifies threads by their creation history. Among other things, this analysis helps to identify which threads are unique and which actions may or may not happen in parallel.

As mentioned, threads are identified by their creation history, which is used as an ID to identify different threads. This history is a sequence of create edges starting with `main`. To prevent such a history to grow to infinity, they define the notion of non-unique thread IDs which may identify multiple threads each.

Formally, the set of possible abstract thread IDs $\mathcal{I}^\#$ is $(\text{main} \cdot \mathcal{P}^*) \times 2^{\mathcal{P}}$, where \mathcal{P} is the set of create edges and \mathcal{P}^* a sequence of such edges. $\langle u, f \rangle \in \mathcal{P}$ refers to an outgoing edge from program point u which creates a thread starting at f . In this notion, IDs of the form $(i, \emptyset) \in \mathcal{I}^\#$ are unique, while $(i, s) \in \mathcal{I}^\#$ are not unique if $s \neq \emptyset$.

As mentioned, these abstract thread IDs are found by a dedicated thread ID analysis. To work correctly, this analysis needs to track the current thread ID as well as a set of create edges already encountered. Thus, the abstract domain for this analysis is $\mathbb{D}_{\text{tID}} = (\mathcal{I}^\# \times 2^{\mathcal{P}})$. An element (id, es) of this domain is the product of the current thread ID id and the set of encountered create edges es .

The main part of the analysis now works as follows: Assume a thread-create edge $(u, \text{create}(f), v)$ is encountered with a state (id, es) . When the id is already non-unique, the new thread is identified with a (possibly new) non-unique ID. In the other case id is unique, i.e., $id = (i, \emptyset)$. Then the analysis checks whether the currently encountered

edge was already encountered before, i.e., if it is present in the set of encountered edges es . If not, the new thread is identified with a new unique ID that is created by appending the current create edge to the first part of $id = (i, \emptyset)$. Otherwise, the new thread is identified with a non-unique ID.

The analysis starts in the main thread with the state $(([main], \emptyset), \emptyset)$. Whenever a thread is created and identified with a new $id \in \mathcal{I}^\#$, this thread is then analyzed with the starting state (id, \emptyset) , beginning at the program point where it starts.

It is worth mentioning, that the $\text{combine}_{\text{tID}}^\#$ function is implemented to ignore the caller state and propagate the callee return state for this analysis.

As of now, the thread ID analysis is always run context-sensitively, i.e., $\text{C}_{\text{tID}} = \text{ID}_{\text{tID}}$ and $\text{context}_{\text{tID}}^\# D = \text{enter}_{\text{tID}}^\# D$. Performing it context-insensitively is not practical. The reason for that is that it leads to cases where a least upper bound of two states (id_1, es_1) and (id_2, es_2) has to be computed. Since abstract IDs from $\mathcal{I}^\#$ are not reasonably comparable, this leads to a notion of "unknown thread ID" or "any possible thread ID" which we want to avoid.

It is much more reasonable to run the thread ID analysis partially context-sensitively, where contexts are differentiated with respect to the current id but not the set of encountered create edges es . Computing the least upper bound of sets of encountered create edges is reasonably possible by taking the union of the sets. With these insights, we can perform the thread ID analysis partially context-sensitively with

$$\text{C}_{\text{tID}} = \mathcal{I}^\# \text{ and } \text{context}_{\text{tID}}^\# (id, es) = id$$

When we now run this partially context-sensitive analysis, it is possible to encounter cases where precision is lost similarly to how we described it in Section 2.6. We show this with the example program from Figure 4.1. In this program, `procedure()` is called two times, once without a thread being created beforehand and once with a thread created. For now, this thread created in Line 14 has the unique ID $([main \cdot \langle u_2, s_{tproc} \rangle])$, assuming the program point before the `create(tproc)` is u_2 and s_{tproc} is the start point of the thread `procedure tproc()`. The two entry states for `procedure()` are $(([main], \emptyset), \emptyset)$ and $(([main], \emptyset), \{\langle u_2, s_{tproc} \rangle\})$, as once the edge $\langle u_2, s_{tproc} \rangle$ was encountered and once it was not. These states both have the same context, i.e., $([main], \emptyset)$, and thus their entry states are joined. The result is the same as the second entry state $(([main], \emptyset), \{\langle u_2, s_{tproc} \rangle\})$. Since we assume for this example, that no thread was created in `procedure()`, the return state is the same as the joined entry state. With the way, the $\text{combine}_{\text{tID}}^\#$ function is defined, this state is used for the point after both calls. In particular, it is used for the state after the call in Line 11. Thus, the `create()` action in Line 14 is actually encountered again with the state $(([main], \emptyset), \{\langle u_2, s_{tproc} \rangle\})$. The result of this is, that the thread created here is now identified with a non-unique

abstract ID ($[main], \{\langle u_2, s_{tproc} \rangle\}$), because this create edge is in the set of encountered create edges.

With this result, some possible race conditions are detected, even though in fact there are none in this program. One reason for that is, that the thread created in Line 14 is no longer identified with a unique abstract ID. Thus, it is no longer possible to know that only one thread runs with the code of procedure $tproc()$ and consequentially, only one thread can perform the assignment instructions in Line 4. Therefore, a data race is found for this assignment.

Furthermore, it is no longer possible to identify, that no thread has been created before the variable assignment in Line 13, because $\langle u_2, s_{tproc} \rangle$ is in the set of encountered create edges of the state at this point. Thus, a possible data race is found also for this assignment instruction.

Note here, that the process we described above is adapted to be better understandable for the reader. Following the actual concepts of static analysis, the corresponding system of constraints for the program points is generated. For this system, a solution can be computed with a fix-point solver.

To combat this loss of precision, we propose a "thread-create" analysis. This analysis checks for each procedure, whether a thread is possibly created between the entry to it and the return. Note that it does not matter if a thread is created in the procedure itself or in another procedure called by the procedure. The domain we use for this analysis is the set of boolean values $\mathbb{D}_{tc} = \{\text{true}, \text{false}\}$. The analysis tracks whether a procedure *may* create a thread. Thus, we encode uncertainty, i.e., "a thread *may* have been created" with true. Similar to the taint analysis from Section 3.1.1, this analysis is context-insensitive by itself and thus, $\mathbb{C}_{tc} = \{\bullet\}$.

In conclusion, the state at some program point $\eta_{tc} [v, \bullet]$ answers the question "May a thread have been created since the entry of the current procedure up to the node v ?".

In the following we give the definitions of the analysis functions for this function:

$$\begin{aligned} \text{init}_{tc}^\# &= \text{false} \\ \llbracket A \rrbracket_{tc}^\# c &= \begin{cases} \text{true} & \text{if } A \equiv (\text{create}(f);) \\ c & \text{else} \end{cases} \\ \text{enter}_{tc}^\# c &= \text{false} \\ \text{combine}_{tc}^\# (c_{cr}, c_{ce}) &= c_{cr} \vee c_{ce} \\ \text{context}_{tc}^\# c &= \bullet \end{aligned}$$

We note that $\text{init}_{tc}^\#$ and $\text{enter}_{tc}^\#$ return false. The reasoning behind this is that when the

```
1  int a;
2
3  void *tproc() {
4    a = 2;
5  }
6  void procedure() {
7    //... (no thread creation)
8  }
9
10 void main() {
11   procedure();
12
13   a = 1;
14   create(tproc);
15
16   procedure();
17 }
```

Figure 4.1: Example program to illustrate the precision loss of partial contexts in the thread ID analysis. A run with the partially context-sensitive thread ID analysis labels the assigns in Line 4 and Line 13 as possible causes for a race condition, even though this program is race free. The fully context-sensitive thread ID analysis correctly finds the absence of data races.

program begins and when a function is entered, the analysis should start with a state describing that no thread has been created. After a call, the state should be true if a thread was created by the caller before or if one was created by the callee. Thus, the $\text{combine}_{tc}^\#$ function performs a logical "or" operation on both states and returns the result.

With this analysis, we can improve the $\text{combine}_{tID}^\#$ of the thread ID analysis. Before, this function was defined as $\text{combine}_{tID}^\# (D_{cr}, D_{ce}) = D_{ce}$. We change this, so this analysis only returns the callee state when a thread was possibly created in the call and otherwise returns the caller state from before:

$$\text{combine}_{tID}^\# (D_{cr}, D_{ce}) = \begin{cases} D_{ce} & \text{if } \eta_{tc} [e_f, c] \\ D_{cr} & \text{else} \end{cases}$$

for an edge $(u, f();, v)$ where c is the corresponding context.

This modification allows the analysis to keep the callee state with the respective set of encountered edges. In the example from Figure 4.1 this means, that after the call in Line 11, the state from before with no encountered edges is kept. This results in no race being detected in Line 13. Furthermore, because the thread created in Line 14 is known to be unique, no race is detected in Line 4.

4.2 Implementation

We briefly described in this section, how we implemented the thread-create analysis from the previous chapter in GOBLINT and how we used it to improve the `threadId` analysis that already exists in the analyzer.

Similar to the `taintPartialContexts` analysis from Section 3.2.1, we implement a new module `threadCreate` that implements the interface seen in Figure 3.2. The implementation is very similar to our formal definition of this analysis from the previous chapter. In the following, we discuss some noteworthy points:

Domain:

The Domain `D` for this analysis only contains the two boolean values, where `false` stands for "definitely no thread was created" and `true` for "maybe a thread was created". For this GOBLINT provides the `MayBool` domain implementing this notion with the corresponding ordering.

Similar to the taint analysis, this analysis does not contribute to the overall context of the analyzer, and thus we chose `Unit` for the module `C`.

Analysis functions:

We handle thread creation in two different functions:

One is the `threadspawn` function, which exists for exactly this purpose. We implement this function to always return `true` for the thread create analysis, as this corresponds to encountering a create edge.

The other function where we handle thread creations is the `special` function. We implement it so that it returns `true` not only when the special function that is called is a thread-creating function, e.g., `pthread_create()`, but also when it is an unknown function. Since we know nothing about unknown functions, we cannot exclude that such a function creates a thread. In any other case besides thread creating and unknown functions, the `special` function propagates the state from before.

All other analysis functions either return the state from before (e.g., `assign` and `return`) or they are implemented exactly as defined in Section 4.1 (e.g. `combine` returns the result of a logical "or" on the caller and callee state).

The query function:

To allow the new thread-create analysis to broadcast its information to other analyses we add a `MayThreadCreate` query. This query is answered by the thread-create analysis with the current state. Other analysis can send a `MayThreadCreate` query to gain the information, whether a thread has been created from the beginning of the current function up to the point where the query is sent. In the case, that the `threadCreate` analysis is not performed, this query is answered with `true` by default, i.e., "A thread *may* have been created".

Improving the `threadId` analysis

In GOBLINT there exists an implementation of the thread ID analysis we introduced in Section 4.1, namely the `threadId` analysis. This analysis was always performed context-sensitively. Thus, we add the option `ana.thread.context.createEdges` to the analyzer. This option is enabled by default, but it allows for removing the set of encountered create edges from the context when it is disabled. We implement this by changing the `context` function of the `threadId` analysis so that it checks this option and returns contexts of the sensitivity specified by it.

To implement our proposed improvements, we change the `combine` function of the `threadId` analysis according to our proposal from Section 4.1. It now sends a `MayThreadCreate` and returns the caller state when the answer is `false`. Otherwise, the callee state is

returned like before our changes.

5 Evaluation

In this chapter, we evaluate our approach. This is split into two parts: First, we test, whether our implementation is sound, i.e., it does not lead to wrong results. After that, we benchmark the implementation on some real-world C programs. The goal of the benchmark is to get a perspective on the precision and computation time of our proposed improvements from Chapter 3 and Chapter 4. Thus, we compare context-insensitive runs with our improvement to context-insensitive ones without the improvement and to context-sensitive runs. In particular, we want to know how much of the precision lost by context-insensitivity can be recovered through our proposed improvements.

5.1 Testing

This section aims to ensure, that the addition of the `taintPartialContexts` analysis to GOBLINT as described in Section 3.2 does not lead to wrong results. Additionally, we want to make sure that just activating the taint analysis on a fully context-sensitive analysis run does not lead to less precise results.

GOBLINT provides an extensive set of regression test cases already. These test edge cases of various features of the analyzer. Each test case comes with a specific configuration, that should be used when executing the test.

To verify our implementation, we ran all regression test cases with their specified configuration but additionally activated the taint analysis. This helps to ensure that no precision is lost, just by activating the taint analysis. These test runs helped to find some bugs, which then were fixed.

For the `threadCreate` analysis we used the same approach, where we ran all regression tests with the thread-create analysis additionally activated.

We also contributed a few new regression tests which specifically test edge cases of the taint and the thread-create analysis. These include tests for each existing analysis, where we included either of our two new analyses, as well as tests for bugs we cleared to show the correct behavior after the fix. The new regression tests also aim to demonstrate how our changes improve the precision of existing analyses.

Additionally, we investigate the results of the benchmarks we describe in the following section. These reinforce the verdict that the taint analysis and the thread-create analysis are sound in their implementation.

5.2 Benchmarking the improved base analysis

In this section, we investigate, if and to what extent the changes we proposed in Section 3.2.2 provide an improvement. In particular, we compare analyses with different configurations in terms of precision and computation time. However, we focus on benchmarking the changes to the base analysis as this is the main variable analysis used in GOBLINT. We do perform benchmarks with the other variable analyses we improved in Section 3.2.2 besides the base analysis.

5.2.1 SV-Benchmarks

The first benchmarking approach we will describe in this section uses the SV-Benchmarks' "Collection of Verification Tasks" [sos]. This collection of verification tasks is "constructed and maintained as a common benchmark for evaluating the effectiveness and efficiency of state-of-the-art verification technology" [sos]. Each verification task consists of a program and a corresponding specification, i.e., a set of properties. The verifier to be benchmarked, i.e., in our case the GOBLINT analyzer with our changes, then performs an analysis run on a given program. After that, it is checked, whether the verifier was able to prove the given set of properties. The properties we focus on for our benchmarks are the following:

- `unreach-call`: A specified function `reach_error()` in the program is never called during runtime.
- `no-overflow`: No integer overflow occurs in the program.
- `no-data-race`: The program contains no race condition.

For each of these properties, we perform three benchmark runs with different configurations on each program for which it is specified that the property holds. The three configurations we use for the benchmark runs are:

- *sens*: The base analysis is performed context-sensitively.
- *insens*: The base analysis is performed context-insensitively.
- *insens taint*: The base analysis is performed context-insensitively with the improvements of the taint analysis.

During each benchmark run with a certain configuration we save different kinds of information to compare later. The most important of these are the computation time per program and the total number of programs, for which the property to prove was proven.

We first investigate the precision. For this, we compare the number of programs for which the property was proven per configuration and property: The raw numbers can be found in Figure 5.2, while the graph in Figure 5.1 shows a visualization of the results.

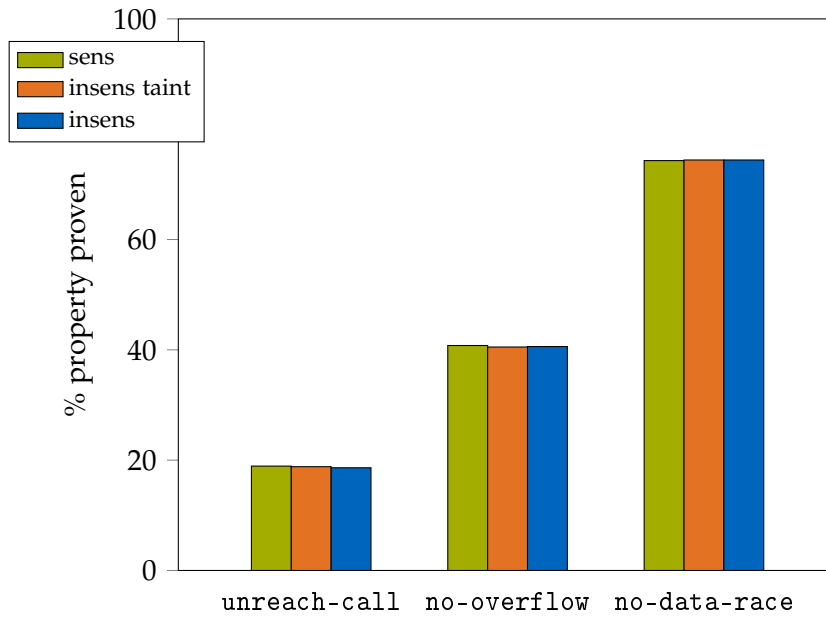


Figure 5.1: Bar graph for the comparison of the three configurations per property. A bar represents the percentage of programs, for which the corresponding property was proven using the corresponding configuration

As we can see in Figure 5.1, there is in fact not a large difference between the three configurations for each property. Our interpretation of this result is, that the choice between our three configurations only has a small effect on the precision of GOBLINT when analyzing real-world programs. In other words, this means, that not a lot of precision is lost by performing the base analysis context-insensitively. The taint analysis can only recover lost precision, so little overall precision loss means, that the taint analysis can only provide little overall improvement.

We now take a closer look at Figure 5.2. We first investigate the results for the

<i>property</i>	<i>config.</i>	<i># proven tasks</i>	<i># total tasks</i>	<i># loss by insens</i>	<i># recovered</i>	<i>% recovered of loss</i>
unreach-call	<i>sens</i>	1925	10173	31	20	64.52
	<i>insens taint</i>	1914				
	<i>insens</i>	1894				
no-overflow	<i>sens</i>	3456	8474	16	-7	-43.75
	<i>insens taint</i>	3433				
	<i>insens</i>	3440				
no-data-race	<i>sens</i>	671	903	-1	0	0.00
	<i>insens taint</i>	672				
	<i>insens</i>	672				

Figure 5.2: Table showing the results of the SV-Benchmarks.

unreach-call property: We see, that with the *insens* configuration, the analyzer was able to prove the property 31 times fewer than with the *sens* configuration, which is less than 0.4% of the total number of tasks. However, it proved the property 20 times more often with the *insens taint* configuration compared to the *insens* configuration. This leads us to the conclusion, that for this property the taint analysis helped to recover about 65% of the precision lost.

As for the results of the no-overflow property, it seems like the usage of the taint analysis leads to less precise results. This property was proven for 7 fewer programs when using the *insens taint* configuration instead of the *insens* configuration. In fact, these results helped us find a bug in our implementation. This bug is now removed and after the fix, the no-overflow property is proven for these 7 programs with the *insens taint* configuration. In conclusion, however, we cannot say that the taint analysis provides an improvement for proving the no-overflow property.

Similarly, the *insens taint* configuration is exactly as precise as the *insens* configuration for the no-data-race property. A timeout is the reason, why the *sens* configuration proved the property for one fewer program than both of the other configurations. Considering this, we conclude, that for this property all three configurations are equally precise.

The table in Figure 5.3 gives information related to the computation times of this benchmark. We first investigate the average computation time of only the tasks, where none of the runs with different configurations terminated with an error. Errors that occurred are mainly timeouts or stack overflow exceptions. A timeout was set to interrupt a task after fifteen minutes of non-termination. We can see, that the differences in the average computation time of tasks without errors are very small. In the case of the

<i>property</i>	<i>config.</i>	<i>avg. cputime all tasks (s)</i>	<i>avg. cputime no errors (s)</i>	<i># Timeouts</i>	<i># Errors</i>
unreach-call	<i>sens</i>	191	30.8	1712	1934
	<i>insens taint</i>	188	31.2	1658	1908
	<i>insens</i>	188	28.1	1563	1799
no-overflow	<i>sens</i>	7.51	0.722	51	130
	<i>insens taint</i>	5.37	0.694	44	45
	<i>insens</i>	5.25	0.627	42	43
no-data-race	<i>sens</i>	1.51	0.510	1	1
	<i>insens taint</i>	0.525	0.525	0	0
	<i>insens</i>	0.517	0.517	0	0

Figure 5.3: Table showing timing results of the SV-Benchmarks. The column *avg. cputime all tasks* contains the average time of all tasks for a property, while *avg. cputime no errors* contains the average time of only the tasks, where none of the three configurations terminated with an error.

unreach-call property, the *insens taint* configuration actually performed worse than the *sens* configuration on average, but not by much. This is the property, where the taint analysis provided the most improvement as we found in the previous paragraph. The most interesting property in this comparison is probably *no-overflow*. Here, both the *insens* and the *insens taint* configuration have a significantly lower average computation time than the *sens* property, when taking all tasks into account. The reason for this is, that the number of errors for each of these two configurations is just about a third when compared to the *sens* configuration. From this, we deduce, that performing the base analysis context-insensitively leads to fewer errors and produces results more quickly compared to performing it context-sensitively. As we found in the previous paragraph, not a lot of precision is lost with the *insens* configuration compared to the *sens* configuration. However, the taint analysis did not provide much improvement for this property.

A similar observation can be made from the timing results of the *no-data-race* condition. For this property, the observation is far less expressive, as only one timeout occurred for the *sens* configuration, compared to no error for the other configurations.

5.2.2 GNU Coreutils

We also benchmark the improved base analysis with another approach. For this, we use modified versions of C programs from the GNU Core Utilities "coreutils" [GNU]. These programs implement "the basic file, shell and text manipulation utilities of the

GNU operating system. These are the core utilities which are expected to exist on every operating system."[GNU]

Combined versions of the coreutil programs are found in a benchmark repository dedicated to providing benchmark programs for the GOBLINT analyzer [Gobb]. A "combined version" of a program is a code file, where all dependencies of included files of the program are merged into one single code file.

For these combined programs we use a feature from the GOBLINT analyzer itself to generate assertions at different points within the program. An assertion is an equality or inequality involving program variables, that holds for every concrete execution of the program. To generate these, the analyzer performs an analysis with a given configuration on a given program. GOBLINT then uses the information it gains to place assertions that it knows are true in the program and produces an output file.

We can then use the resulting file with generated assertions to compare other analysis runs with different configurations. The metric for these comparisons is the number of proven assertions. The configuration which is used to generate the assertions must be at least as precise as the most precise configuration of the runs we want to compare. This benchmarking approach can produce meaningful results only under these conditions. For this reason, we generate assertions with a configuration that performs the base analysis context-sensitively. To do this we follow an approach similar to Julian Erhard, where he generated assertions for these programs context-insensitively. A detailed description of his approach can be found at [Erh].

We compare the same configurations we compared in Section 5.2.1.

The results of the benchmark runs on the coreutil programs can be seen in the table in Figure 5.4. We omit the results for two further programs `cp_comb.c` and `mv_comb.c` because for these the analyzer proved some assertions to evaluate to "false" with the *insens* and *insens taint*. This hints at bugs in the analyzer which are unrelated to our changes.

The number of total assertions differs between the *sens* and both *insens* (*taint*) configurations for some programs. The reason for this is, that an assertion in a function can be checked multiple times, once for each context the function is evaluated with. Therefore, we compare percentages of proven assertions for this benchmark.

We make the following observations: The percentage of the precision loss that is recovered reaches over 16% for one program while it stays below 2% for all other programs. For three of the eight programs shown, no improvement was achieved at all. From this, we conclude that how much the taint analysis can improve the precision loss depends a lot on the program analyzed.

We also note that the "% recovered of total" is rather low. That is the percentage of assertions, which were found with the *insens taint* but not the *insens* configuration,

<i>program</i>	<i>config.</i>	<i># asserts proven</i>	<i># total asserts</i>	<i>% loss by insens</i>	<i>% recovered of total</i>	<i>% recovered of loss</i>
cksum_comb.c	<i>sens</i>	1998	1998	1.30	0.00	0.00
	<i>insens taint</i>	1972	1998			
	<i>insens</i>	1972	1998			
cut_comb.c	<i>sens</i>	3992	3992	0.30	0.00	0.00
	<i>insens taint</i>	3979	3992			
	<i>insens</i>	3979	3992			
dd_comb.c	<i>sens</i>	4462	4464	2.97	0.49	16.54
	<i>insens taint</i>	4359	4472			
	<i>insens</i>	4337	4472			
df_comb.c	<i>sens</i>	8834	8834	12.68	0.05	0.36
	<i>insens taint</i>	7718	8834			
	<i>insens</i>	7714	8834			
du_comb.c	<i>sens</i>	9810	9810	3.41	0.05	1.50
	<i>insens taint</i>	9469	9798			
	<i>insens</i>	9464	9798			
nohup_comb.c	<i>sens</i>	3397	3397	0.77	0.00	0.00
	<i>insens taint</i>	3371	3397			
	<i>insens</i>	3371	3397			
ptx_comb.c	<i>sens</i>	5786	5786	4.46	0.07	1.55
	<i>insens taint</i>	5532	5786			
	<i>insens</i>	5528	5786			
tail_comb.c	<i>sens</i>	4806	4806	0.50	0.00	0.00
	<i>insens taint</i>	4782	4806			
	<i>insens</i>	4782	4806			

Figure 5.4: Table showing the results of the coreutils with generated assertions

never surpasses 0.5%. Therefore, we conclude that the effect of our changes on the overall precision of the analyzer is very limited.

5.3 Benchmarking the improved `threadId` analysis

We only performed a partial benchmark for our improvement of the `threadId` analysis using the thread-create analysis. For this, we again used the GNU coreutil programs, which we described in Section 5.2.2. This time, however, instead of comparing the number of correctly proven assertions, we compare the number of race conditions, the GOBLINT analyzer finds in each program with each configuration. We compared these configurations:

- *sens*: The `threadId` analysis is performed context-sensitively.
- *part-sens*: The `threadId` analysis is performed partially context-sensitively, i.e., only sensitive with respect to the thread ID but not with respect to the set of encountered create edges. Concretely, the option `ana.thread.context.createEdges` we added is disabled in this configuration.
- *part-sens taint*: Same configuration like *part-sens* but with the improvements of the thread-create analysis enabled.

The results of this benchmark are as follows: For each of the coreutil programs, the analyzer finds between 5 and 8 race conditions. However, there is no difference between the three configurations. This means that for each program, the number of race conditions found is the same regardless of the configuration chosen. This means that with this benchmarking setup, we do not record any precision loss that results from partial context-sensitivity. Since no precision is lost, we cannot see, if the thread-create analysis provides an improvement.

In conclusion, there definitely exist cases, where running the `threadId` partially context-sensitive leads to precision loss, which the thread-create analysis can reduce. We have shown this by adding a regression test of one such case to the GOBLINT regression tests, similar to the example in Figure 4.1.

However, such cases may only exist very rarely in real-world programs.

It has to be noted that our benchmark for this analysis was only very small in scale. More extensive benchmarks are necessary to make a more meaningful statement. The next step in the future is to perform a benchmark run of the SV-Benchmarks' `unreach-call` property as described in Section 5.2.1.

6 Conclusions

In this thesis, we introduced static analysis with systems of constraints. After expanding the approach to interprocedural analysis with various kinds of context-sensitivity, we identified a source of precision loss that can occur with context-insensitive and partially context-sensitive analyses. In the following two chapters, we focused on two different kinds of analyses, where this loss of precision occurs and proposed an improvement that reduces the precision lost for each kind. We implemented both of our proposed approaches in the GOBLINT analyzer. Afterward, we tested and benchmarked our implementation.

The following conclusions originate mostly from the results of the benchmarks we performed with the base analysis of GOBLINT. This analysis is based on the values-of-variables analysis and expands on it. Thus, we suspect our conclusions to hold for this type of analysis in general.

In contrast to our expectation, the precision loss of performing the base analysis context-insensitively instead of context-sensitively does not lead to a great loss of precision in general.

Additionally, the benefit of a faster computation time through context-insensitivity is also smaller than expected for most programs. There are however certain programs, where a context-sensitive base analysis struggles as it terminates with a timeout or stack overflow, while a context-insensitive base analysis terminates much quicker without an error. A shallow investigation of this phenomenon showed, that a large portion of the programs for which these errors occurred are ones with recursive functions. We note here, that this phenomenon was acknowledged by Smaragdakis et al. [SKB14], where they state that it is an often reported problem of context-sensitivity, that it scales rather well in some cases, while in other cases it becomes orders-of-magnitudes more expensive.

Adding the taint analysis we proposed in this thesis only provided a limited improvement in precision to the context-insensitive case. However, as mentioned above only a small portion of precision was lost through context-insensitivity. This limits the improvement, which the taint analysis can provide as its addition cannot make a context-insensitive analysis more precise than a context-sensitive one. Comparing the amount of precision lost through context-insensitivity with the amount of precision

recovered by the taint analysis makes the benefit of the taint analysis much more noticeable.

For computation time we note that a context-insensitive base analysis improved with the taint analysis takes about as long as a context-sensitive base analysis with some variance, provided neither of them terminates with an error. However, just like the context-sensitive analysis by itself, the improved analysis terminates with an error in much fewer cases than the context-sensitive analysis.

Our final note on the combination of a context-insensitive base analysis with the improvement of the taint analysis is the following: This configuration has its best use in the case when errors such as timeouts and stack overflows should be avoided as much as possible, but precision is more important than computation time. In this case, errors are avoided through context-insensitivity and precision is maximized through the taint analysis.

The benchmark for our changes to the thread ID analysis was extremely limited, and thus we cannot give more than a first impression based on its results: We did not record any precision loss. Thus, we suspect that for most real-world programs performing the thread ID analysis partially context-sensitively, as introduced in this thesis, does not result in a large precision loss if any. We cannot comment on the performance of the thread create analysis. The reason for this is, that this analysis can only reduce existing precision loss, but we recorded none in the benchmark. Further benchmarks are necessary for a more expressive conclusion on this kind of analysis.

We would like to perform further benchmarks on the thread ID and thread create analysis. A run of the SV-Benchmarks' tasks with the `no-data-race` property is a suitable next step.

Furthermore, we want to inspect the benchmark results for the base analysis more closely. We hope to find some sort of characterization of programs, where the taint analysis provided an improvement.

Lastly, we note that similar approaches to the ones we proposed in the taint and thread create analysis can be applied to other kinds of analyses. In general, if it can be identified that some part of a state was not altered during a procedure call, that part can be kept from the caller state and does not have to be overwritten by potentially less precise information from the callee state. We would like to experiment with applying this principle to other analyses in future work.

List of Figures

1.1	C code sample with multiple function calls to illustrate the precision loss of partial context-sensitivity.	2
2.1	Example program (left) and corresponding CFG (right).	6
2.2	Example program (left) and corresponding CFGs for main (middle) and incr (right). The asserts in Line 14 and Line 15 are proven with a context-sensitive values-of-variables analysis. However, a context-insensitive values-of-variables analysis fails to prove both.	11
3.1	Schematic directory structure of GOBLINT. Adapted from [Api14]. . . .	19
3.2	Simplified Interface for implementing analyses in GOBLINT	21
3.3	Implementation of the helper <code>taint_lval</code> and the <code>assign</code> function for the taint analysis in GOBLINT.	23
4.1	Example program to illustrate the precision loss of partial contexts in the thread ID analysis. A run with the partially context-sensitive thread ID analysis labels the assigns in Line 4 and Line 13 as possible causes for a race condition, even though this program is race free. The fully context-sensitive thread ID analysis correctly finds the absence of data races.	35
5.1	Bar graph for the comparison of the three configurations per property. A bar represents the percentage of programs, for which the corresponding property was proven using the corresponding configuration	41
5.2	Table showing the results of the SV-Benchmarks.	42
5.3	Table showing timing results of the SV-Benchmarks. The column <i>avg. cputime all tasks</i> contains the average time of all tasks for a property, while <i>avg. cputime no errors</i> contains the average time of only the tasks, where none of the three configurations terminated with an error.	43
5.4	Table showing the results of the coreutils with generated assertions . .	45

Bibliography

- [Api14] K. Apinis. “Frameworks for analyzing multi-threaded C.” PhD thesis. Technische Universität München, 2014.
- [ASV12] K. Apinis, H. Seidl, and V. Vojdani. “Side-effecting constraint systems: a swiss army knife for program analysis.” In: *Asian Symposium on Programming Languages and Systems*. Springer. 2012, pp. 157–172.
- [Cil] Cil. *C Intermediate Language (CIL)*. <https://github.com/goblint/cil>. Last Accessed: 2023-02-01. Goblint.
- [Cou21] P. Cousot. *Principles of Abstract Interpretation*. MIT Press, 2021.
- [Erh] J. Erhard. *sv-benchmarks/c/goblint-coreutils/README.txt*. <https://gitlab.com/sosy-lab/benchmarking/sv-benchmarks/-/blob/9930354e/c/goblint-coreutils/README.txt>. Last Accessed: 2023-02-01.
- [GNU] GNU. *Coreutils - GNU core utilities*. <https://www.gnu.org/software/coreutils/>. Last Accessed: 2023-02-01. GNU Operating System.
- [Goba] Goblint. *Goblint - A static analyzer for multi-threaded C programs, specializing in finding concurrency bugs*. <https://goblint.in.tum.de/home>. Last Accessed: 2023-02-01. Goblint.
- [Gobb] Goblint. *Goblint benchmark suite*. <https://github.com/goblint/benchmark/tree/master>. Last Accessed: 2023-02-01. Goblint.
- [Li+18] Y. Li, T. Tan, A. Møller, and Y. Smaragdakis. “Scalability-first pointer analysis with self-tuning context-sensitivity.” In: *Proceedings of the 2018 26th ACM joint meeting on european software engineering conference and symposium on the foundations of software engineering*. 2018, pp. 129–140.
- [Li+20] Y. Li, T. Tan, A. Møller, and Y. Smaragdakis. “A principled approach to selective context sensitivity for pointer analysis.” In: *ACM Transactions on Programming Languages and Systems (TOPLAS)* 42.2 (2020), pp. 1–40.
- [LX19] J. Lu and J. Xue. “Precision-preserving yet fast object-sensitive pointer analysis with partial context sensitivity.” In: *Proceedings of the ACM on Programming Languages* 3.OOPSLA (2019), pp. 1–29.

- [Oh+14] H. Oh, W. Lee, K. Heo, H. Yang, and K. Yi. “Selective context-sensitivity guided by impact pre-analysis.” In: *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design and Implementation*. 2014, pp. 475–484.
- [RY20] X. Rival and K. Yi. *Introduction to static analysis: an abstract interpretation perspective*. Mit Press, 2020.
- [Sch+23] M. Schwarz, S. Saan, H. Seidl, J. Erhard, and V. Vojdani. “Clustered Relational Thread-Modular Abstract Interpretation with Local Traces.” In: *arXiv preprint arXiv:2301.06439* (2023).
- [SKB14] Y. Smaragdakis, G. Kastrinis, and G. Balatsouras. “Introspective analysis: context-sensitivity, across the board.” In: *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design and Implementation*. 2014, pp. 485–495.
- [sos] sosy-lab. *SV-Benchmarks*. <https://gitlab.com/sosy-lab/benchmarking/sv-benchmarks>. Last Accessed: 2023-02-01. sosy-lab.
- [TN19] M. Thakur and V. K. Nandivada. “Compare less, defer more: Scaling value-contexts based whole-program heap analyses.” In: *Proceedings of the 28th International Conference on Compiler Construction*. 2019, pp. 135–146.